

ABY³: A MIXED PROTOCOL FRAMEWORK FOR MACHINE LEARNING

Peter Rindal
Payman Mohassel

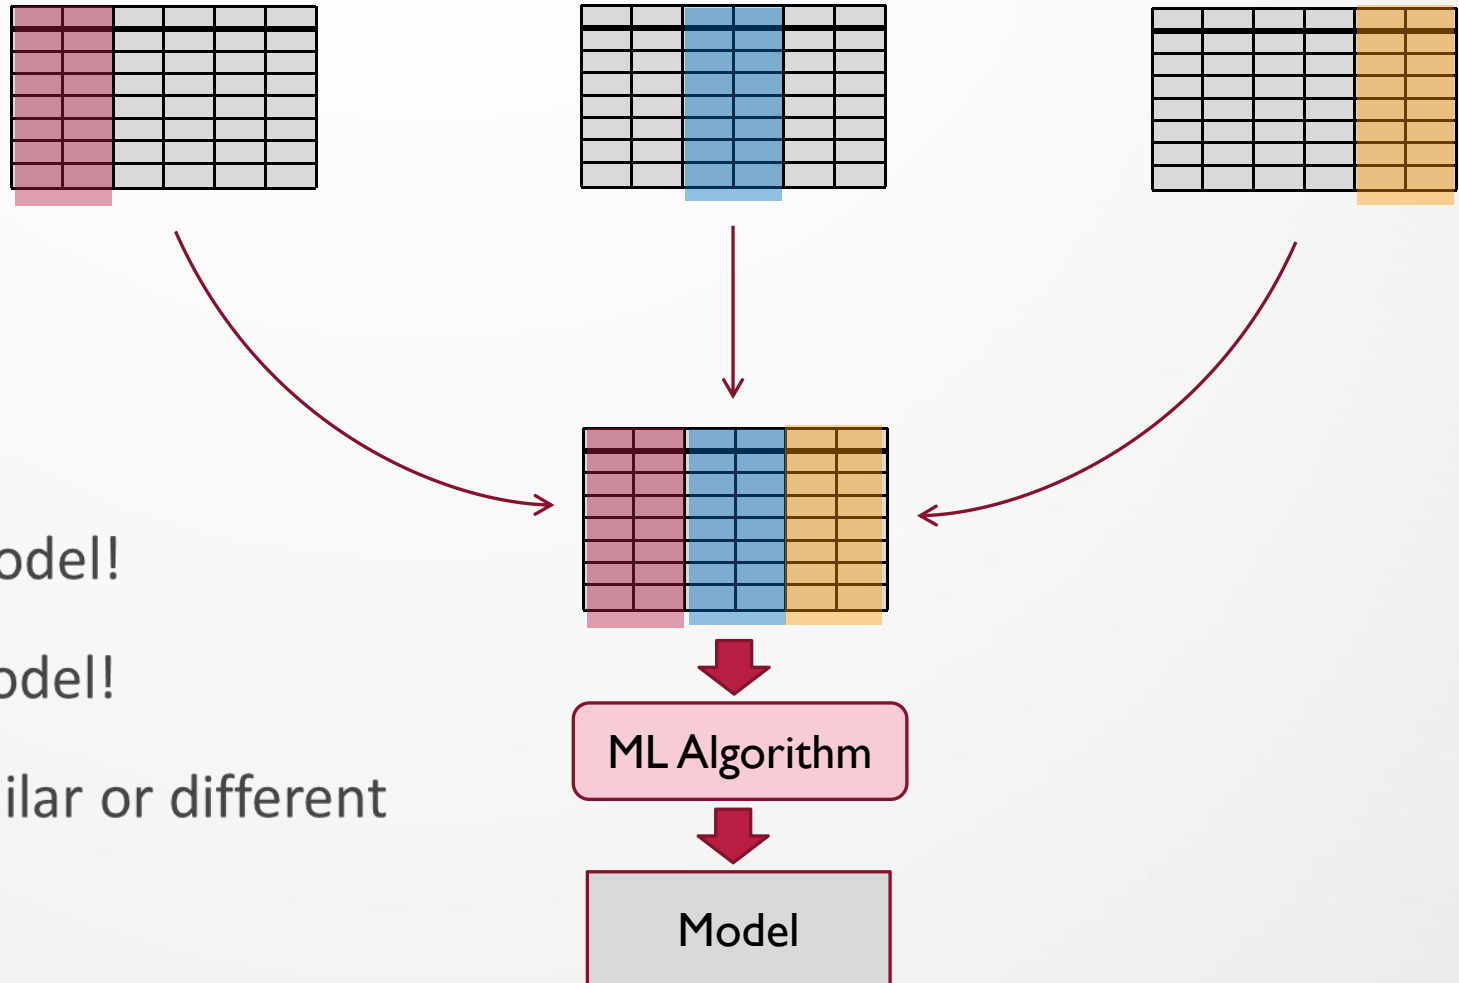
VISA
Research

Machine Learning and Privacy

- **Cognitive tasks:** voice, facial recognition
- **Medical:** genetic testing, disease prediction
- **Financial:** fraud detection, credit rating
- Inference:
 - User does not want to share their data
 - Model owner does not want to share model
- Training:
 - Requires large data sets, often from different sources

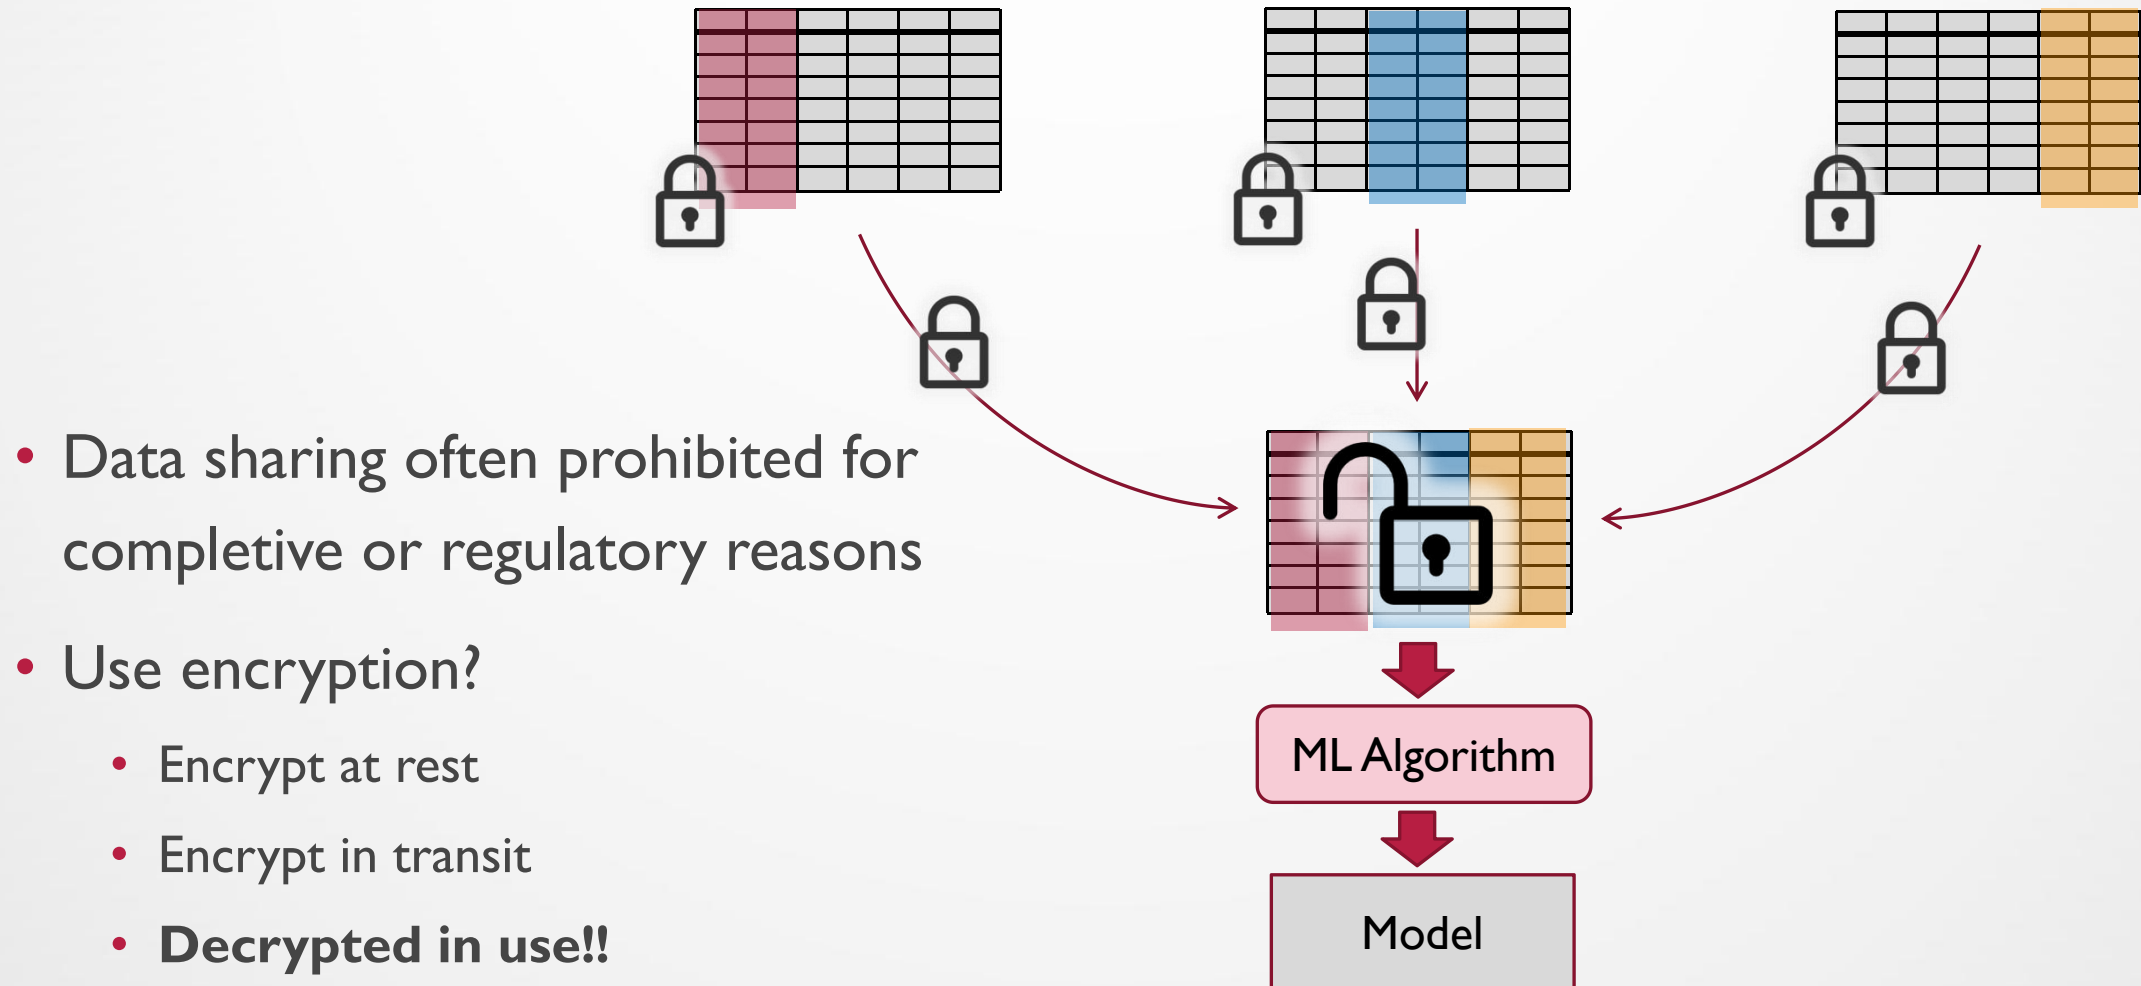


Training A better Model



- More data → better model!
- Richer data → better model!
- Each party can have similar or different types of data

Security Concerns



Our Results

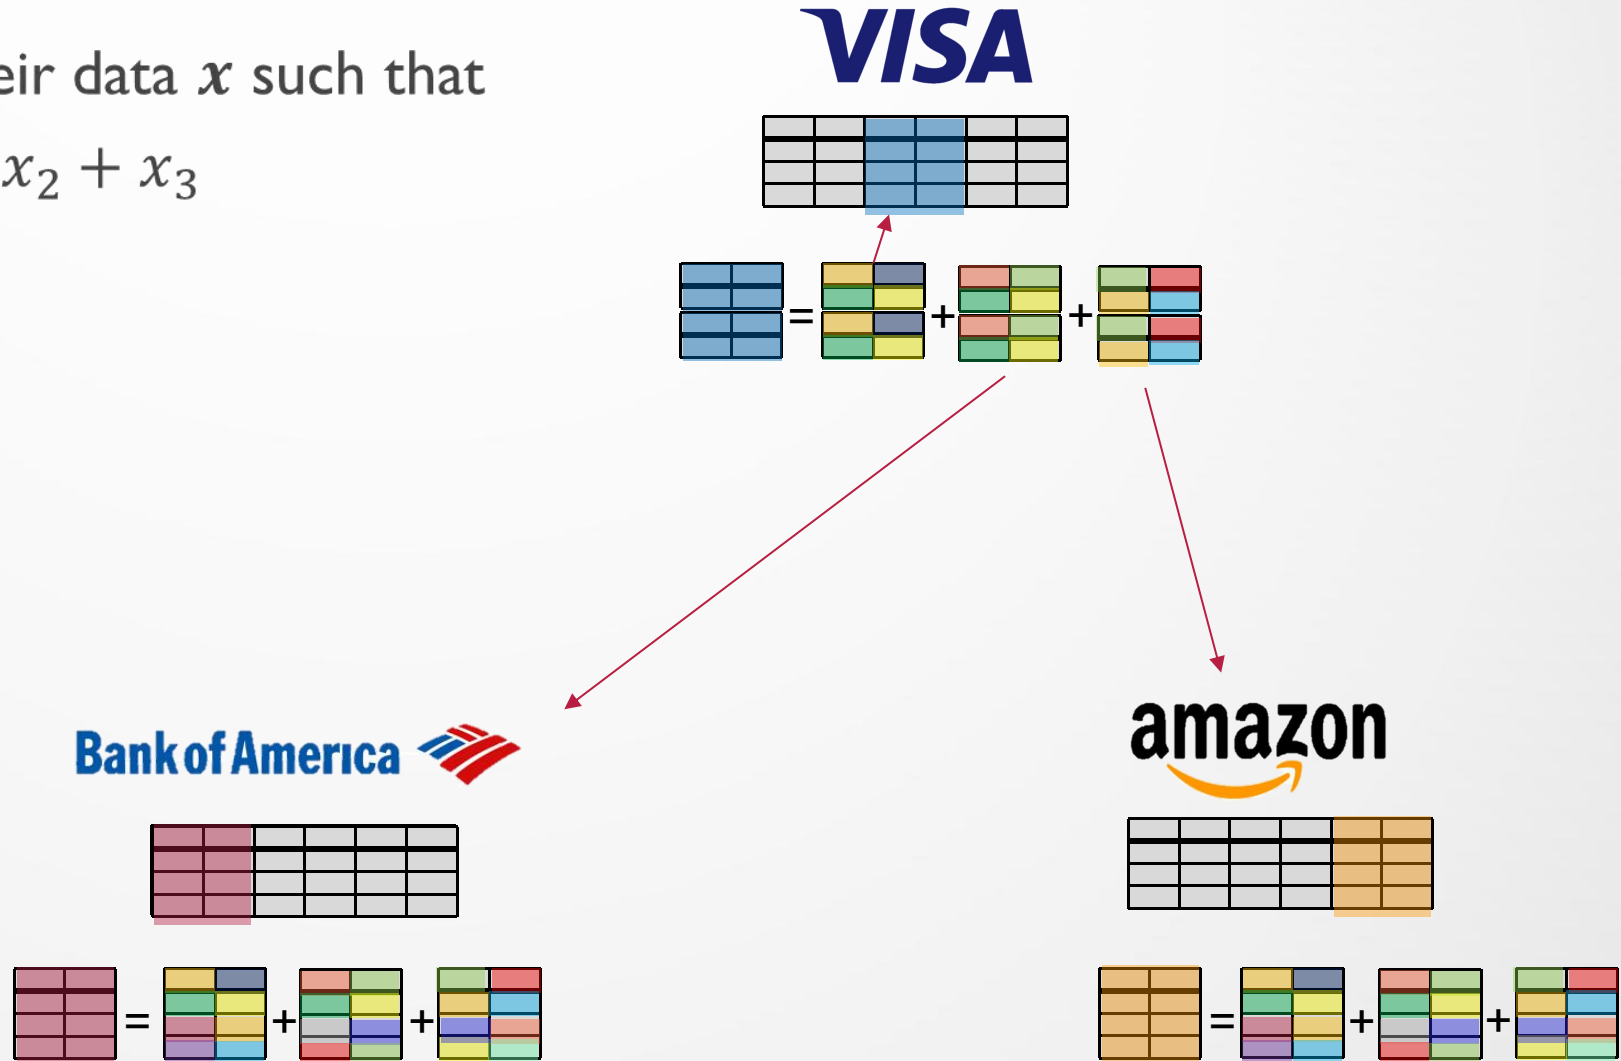
- New three party MPC Protocols:
 - Efficient support for fixed-point arithmetics
 - Improved matrix multiplication
 - Efficient piece-wise polynomial evaluation
 - Conversions between **A**rithmetic, **B**oolean, and **Y**ao secret shares
- Always encrypted machine learning **training and inference**:
 - Linear Regression
 - Logistic Regression
 - Neural Networks
 - Extendable to other models

Protocols and Building block

Always Encrypted

- Each party “Encrypts” their data x such that

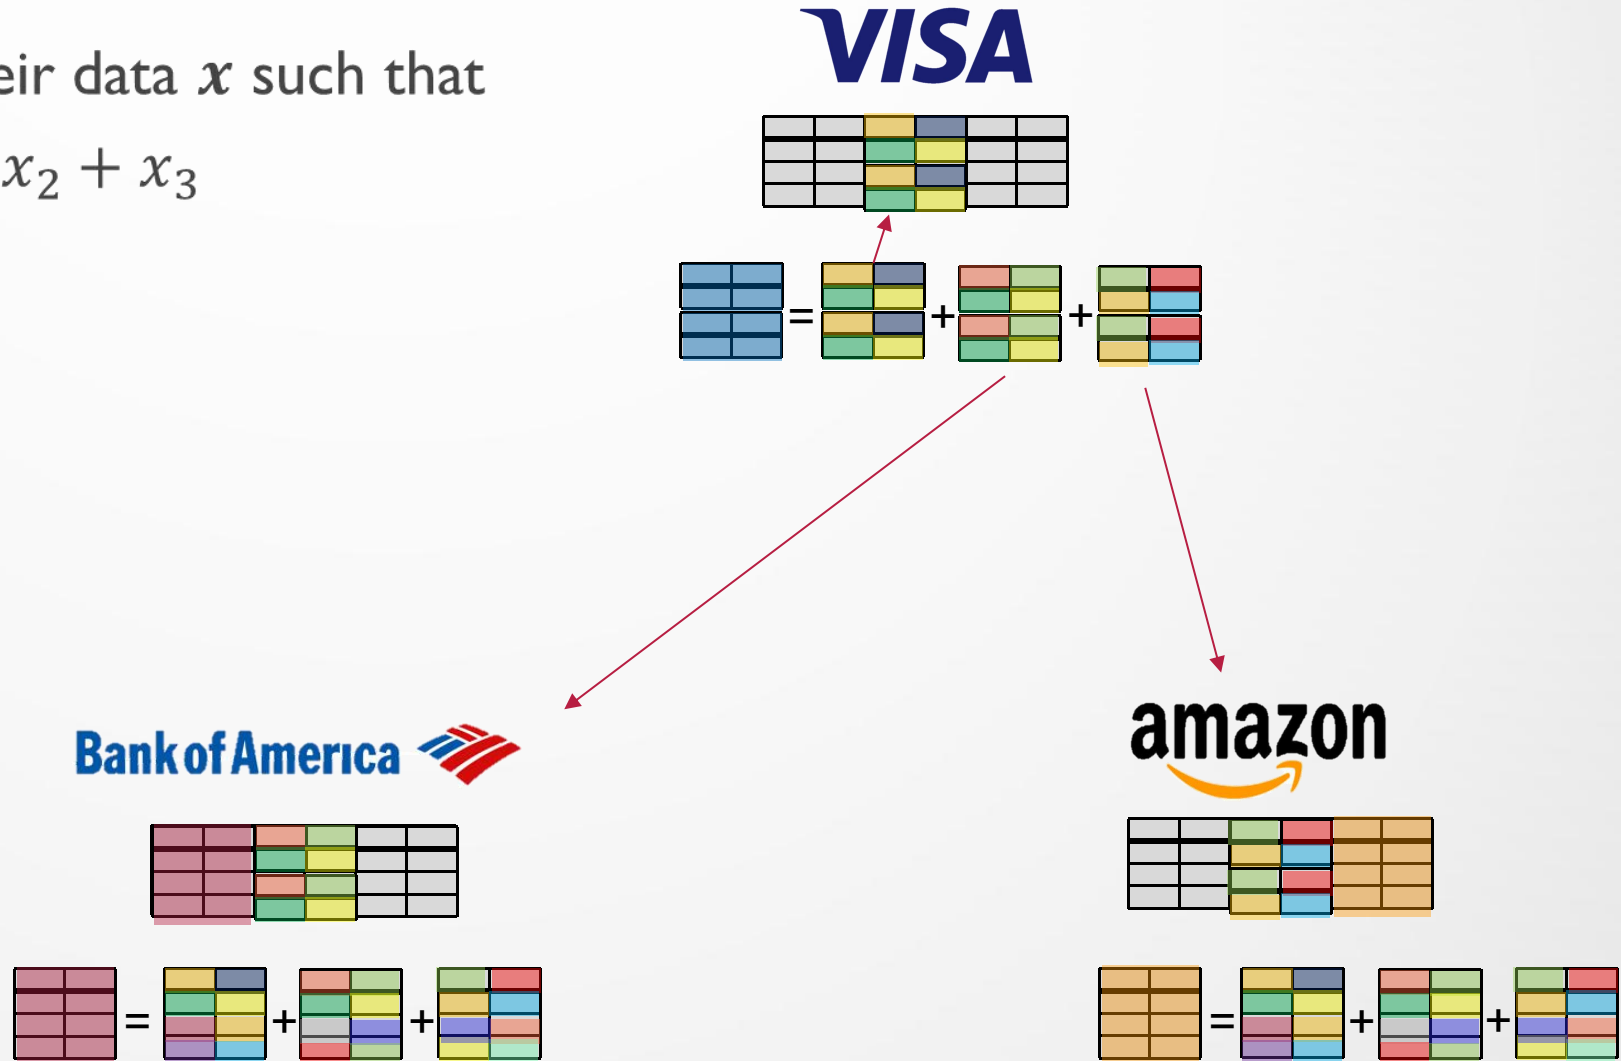
$$x = x_1 + x_2 + x_3$$



Always Encrypted

- Each party “Encrypts” their data x such that

$$x = x_1 + x_2 + x_3$$



Always Encrypted


- Each party “Encrypts” their data x such that

$$x = x_1 + x_2 + x_3$$


- Share Types:

- Arithmetic:** $\llbracket x \rrbracket^A \Rightarrow x_1 + x_2 + x_3$
- Boolean:** $\llbracket x \rrbracket^B \Rightarrow x_1 \oplus x_2 \oplus x_3$
- Yao G.C.:** $\llbracket x \rrbracket^Y \Rightarrow LSB(x_1 \oplus x_2)$


Bank of America



VISA



amazon

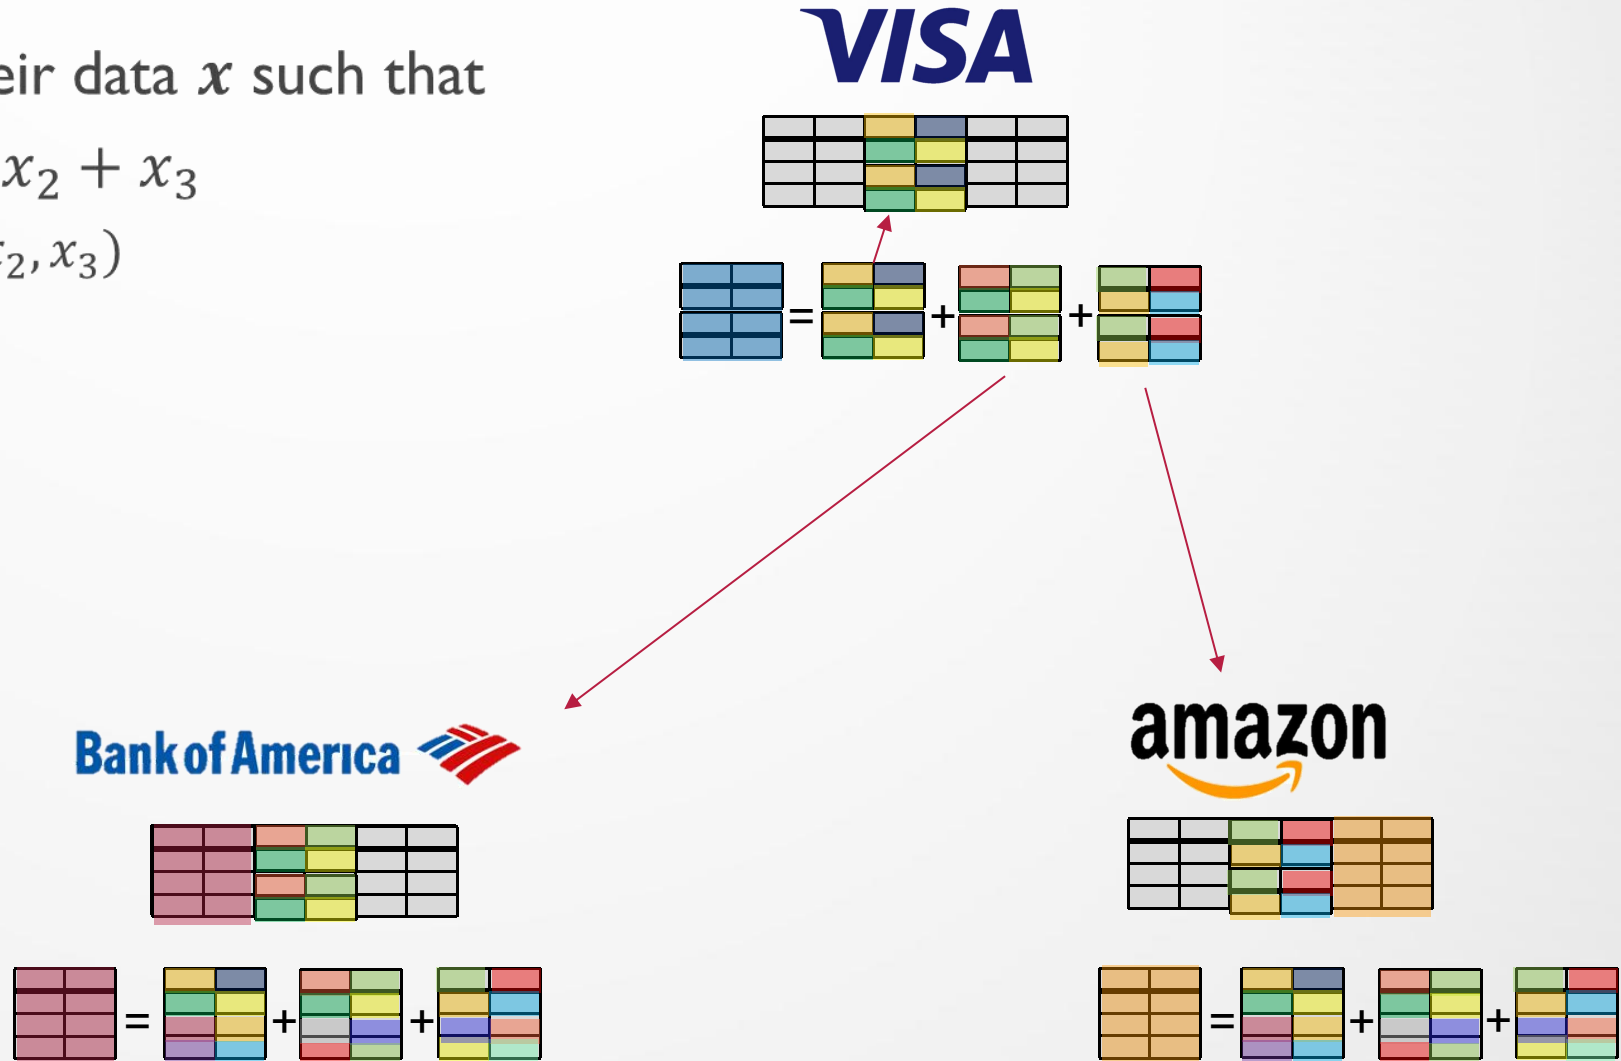


Always Encrypted

- Each party “Encrypts” their data x such that

$$x = x_1 + x_2 + x_3$$

$$\llbracket x \rrbracket := (x_1, x_2, x_3)$$



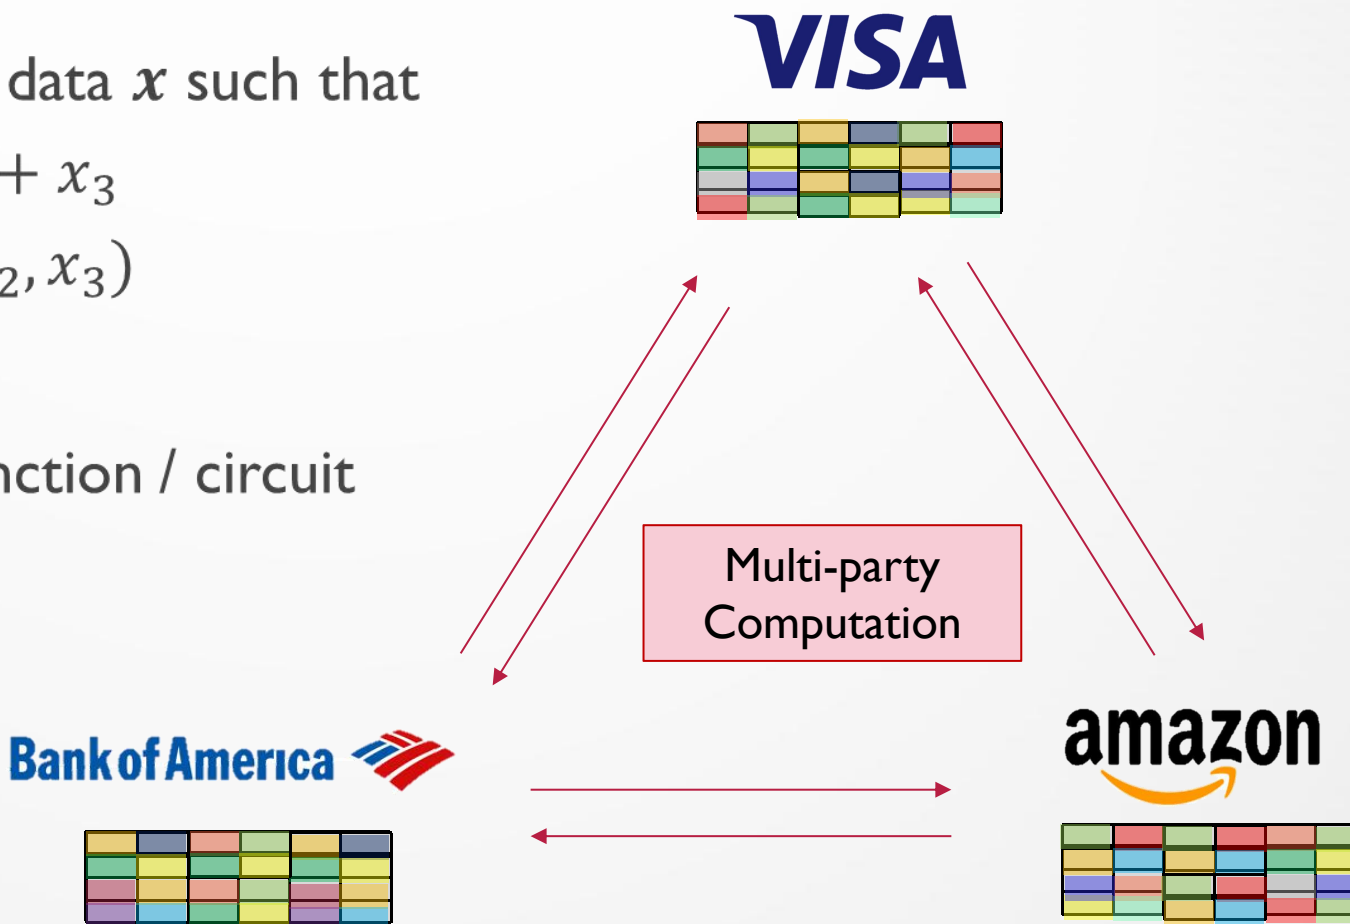
Always Encrypted

- Each party “Encrypts” their data x such that

$$x = x_1 + x_2 + x_3$$

$$\llbracket x \rrbracket := (x_1, x_2, x_3)$$

- Possible to compute any function / circuit in this format



How to Compute on Shared Data

- Say we have shared $\llbracket x \rrbracket, \llbracket y \rrbracket$

VISA

x_1, y_1

- Addition $\llbracket z \rrbracket = \llbracket x \rrbracket + \llbracket y \rrbracket$:

$$z_i = x_i + y_i$$

x_2, y_2

Bank of America 

x_3, y_3

amazon 

How to Compute on Shared Data

- Multiplication $[[z]] = [[x]] * [[y]]$:
 - $z_i \neq x_i * y_i$
 - Use “replicated secret sharing”

VISA

x_1, y_1

x_2, y_2

Bank of America 

x_3, y_3

amazon 

How to Compute on Shared Data

- Multiplication $[[z]] = [[x]] * [[y]]$:

- $z_i \neq x_i * y_i$
- Use “replicated secret sharing”
- Observe:

$$xy = (x_1 + x_2 + x_3)(y_1 + y_2 + y_3)$$

$$\begin{aligned}
 &= x_1y_1 + x_1y_2 + x_1y_3 \\
 &\quad + x_2y_1 + x_2y_2 + x_2y_3 \\
 &\quad + x_3y_1 + x_3y_2 + x_3y_3
 \end{aligned}$$

- Let

$$z_1 = x_1y_1 + x_1y_3 + x_3y_1$$

$$z_2 = x_1y_2 + x_2y_2 + x_2y_3$$

$$z_3 = x_2y_3 + x_3y_2 + x_3y_3$$

VISA

x_1, y_1
 x_3, y_3

x_1, y_1
 x_2, y_2

x_2, y_2
 x_3, y_3

Bank of America 

amazon 

How to Compute on Shared Data

- Multiplication $[[z]] = [[x]] * [[y]]$:

- $z_i \neq x_i * y_i$
- Use “replicated secret sharing”
- Observe:

$$xy = (x_1 + x_2 + x_3)(y_1 + y_2 + y_3)$$

$$= \begin{matrix} x_1y_1 & x_1y_2 & x_1y_3 \\ +x_2y_1 & +x_2y_2 & +x_2y_3 \\ +x_3y_1 & +x_3y_2 & +x_3y_3 \end{matrix}$$

- Let

$$z_1 = x_1y_1 + x_1y_3 + x_3y_1$$

$$z_2 = x_1y_2 + x_2y_2 + x_2y_2$$

$$z_3 = x_2y_3 + x_3y_2 + x_3y_3$$

VISA

x_1, y_1, z_1
 $x_3, y_3,$

$x_1, y_1,$
 x_2, y_2, z_2


$x_2, y_2,$
 x_3, y_3, z_3

Bank of America 

amazon 


Decimal Multiplications in Integer Group

- Treat bottom bits as factional bits
- Multiplication drops top/bottom d bits



Decimal Multiplications in Integer Group

- Treat bottom bits as fractional bits
- Multiplication drops top/bottom d bits
- Mohassel Zhang '17
 - Introduced secret shared version
 - Adds small rounding errors
 - Only works on **2-out-of-2 secret sharing**
- We extend to **any honest majority**
 - Briefly switch to 2-out-of-2 sharing




Preprocess: $\llbracket r \rrbracket \leftarrow \mathbb{Z}_{2^m}, \quad \llbracket r' \rrbracket = \llbracket r \rrbracket / 2^\ell$

1. $\llbracket z' \rrbracket = \llbracket x \rrbracket * \llbracket y \rrbracket$
2. $t = \text{Reveal}(\llbracket z' \rrbracket - \llbracket r \rrbracket)$
3. $\llbracket z \rrbracket = t / 2^\ell + \llbracket r' \rrbracket$

Decimal Multiplications in Integer Group

- Treat bottom bits as factional bits
- Multiplication drops top/bottom d bits
- Mohassel Zhang '17
 - Introduced secret shared version
 - Adds small rounding errors
 - Only works on **2-out-of-2 secret sharing**
- We extend to **any honest majority**
 - Briefly switch to 2-out-of-2 sharing
 - Round preserving
 - Malicious secure



Preprocess: $\llbracket r \rrbracket \leftarrow \mathbb{Z}_{2^m}, \quad \llbracket r' \rrbracket = \llbracket r \rrbracket / 2^\ell$

1. $\llbracket z' \rrbracket = \llbracket x \rrbracket * \llbracket y \rrbracket$
2. $t = \text{Reveal}(\llbracket z' \rrbracket - \llbracket r \rrbracket)$
3. $\llbracket z \rrbracket = t/2^\ell + \llbracket r' \rrbracket$

Share #1

Share #2

Matrix Multiplication on Shared Data

- Given, $X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nd} \end{bmatrix}, Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_d \end{bmatrix}$

- Compute $Z = X \times Y$

- Classic “rows \times columns” algorithm. E.g. $z_i = \sum_j x_{ij} y_j$

- Each multiplication requires communication

- Communication = nd numbers!

- Recall “half” a multiplication is done locally.

- Do half multiplies for $z_i = \sum_j x_{ij} y_j$

- Only send final result z_i

- d times less communication

$$z_1 = x_1 y_1 + x_1 y_3 + x_3 y_1$$

$$z_2 = x_1 y_2 + x_2 y_2 + x_2 y_2$$

$$z_3 = x_2 y_3 + x_3 y_2 + x_3 y_3$$

VISA

$$\begin{matrix} x_1, y_1, z_1 \\ x_3, y_3, z_3 \end{matrix}$$

$$\begin{matrix} x_1, y_1, \\ x_2, y_2, z_2 \end{matrix}$$

$$\begin{matrix} x_2, y_2, \\ x_3, y_3, z_3 \end{matrix}$$

Piece-wise Polynomial

- Let

- $t_0 = -\infty$
- $t_1 = \text{first threshold}$
- ...
- $t_m = \infty$

- Strategy:


$$f(\llbracket x \rrbracket) = \sum_{i=0}^{m-1} \overbrace{\left(\llbracket x \rrbracket > t_i \wedge \llbracket x \rrbracket \leq t_{i+1} \right)}^{\text{range test}} * f_i(\llbracket x \rrbracket)$$

- Polynomial $f_i(\cdot)$ easy to compute

- How to compute $c_i = \llbracket x \rrbracket > t_i \wedge \llbracket x \rrbracket \leq t_{i+1}$?


- $\llbracket t_i - x \rrbracket^B = \text{Convert-to-binary}(\llbracket t_i - x \rrbracket)$
- $\llbracket b_i \rrbracket^B = \text{Most-significant-bit}(\llbracket t_i - x \rrbracket^B) \quad // \quad b_i = (x > t_i)$
- $\llbracket c_i \rrbracket = \text{Converty-to-arithmetic}(\llbracket b_i \rrbracket \wedge \neg \llbracket b_{i+1} \rrbracket)$

- Total cost: $O(\log m)$ rounds and $O(m)$ bits of communication



Conversion

- Need to convert secret sharing of $x \in \mathbb{Z}_{2^k}$:
 - Arithmetic to binary: $[[x]]^A \rightarrow [[\vec{x}]]^B = [[x_1]]^B, \dots, [[x_k]]^B$
 - Binary to Arithmetic: $[[\vec{x}]]^B \rightarrow [[x]]^A$
- Strategy:
 - $[[x]]^A = (x_1, x_2), (x_2, x_3), (x_3, x_1)$
 - Parties provide x_1, x_2, x_3 as input to the binary MPC protocol
 - Evaluate the circuit: $[[\vec{x}]]^B = [[\vec{x}_1]]^B + [[\vec{x}_2]]^B + [[\vec{x}_3]]^B$
- Optimizations:
 - Implement + using low depth circuit, $depth = \log k$
 - Have party 1 locally compute and input $x_1 + x_2$
- Also implement conversion to/from garbled circuits
 - Garbled circuits allow low latency for high depth circuit



Malicious Security

- We extend our semi-honest protocol to the malicious setting
- If any **one** party tries to cheat, the protocol can detect it.
 - Leverage replicated secret sharing
- Most operations roughly 4 times more expensive

VISA

x_1, y_1, z_1
 x_3, y_3, z_3


x_1, y_1, z_1
 x_2, y_2, z_2

x_2, y_2, z_2
 x_3, y_3, z_3

Application: Machine Learning

Linear Regression on Shared Data

- Given, $X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nd} \end{bmatrix}$, $y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$
- Find a linear function f s.t. $f(X_i) \approx y_i$



Linear Regression on Shared Data


- Given, $X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nd} \end{bmatrix}, y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$

- Find a linear function f s.t. $f(X_i) \approx y_i$

- Model: $f_w(X_i) = \sum_{j=1}^d x_{ij} w_j$

- Update function:

$$w_j = w_j - \alpha(X_i w - y_i)x_{ij}$$



- Cost function: $C(w) = \frac{1}{2} (f_w(X) - y)^2$
(a.k.a. L2 norm)

Linear Regression on Shared Data

- Given, $X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nd} \end{bmatrix}, y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$

- Find a linear function f s.t. $f(X_i) \approx y_i$

- Model: f_w

- Update function:

$$w_j = w_j - \alpha (X_i w - y_i) x_{ij}$$


Learning
Rate

Error
Magnitude

Error
Direction

Cost function: $C(w) = \frac{1}{2} (f_w(X) - y)^2$

(a.k.a. L2 norm)



Batch Linear Regression on Shared Data


- Given, $X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nd} \end{bmatrix}$, $Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$

Current Batch

- Find a linear function f s.t. $f(X_i) \approx y_i$
- Update in batches:
 - For $B \subset \{1, 2, \dots, n\}$

$$\mathbf{w} = \mathbf{w} - \frac{1}{|B|} \alpha X_B^T \times (X_B \times \mathbf{w} - Y_B)$$

- Must do many matrix multiplications



Logistic Regression

- Instead of computing



$$\mathbf{w} = \mathbf{w} - \frac{1}{|B|} \alpha X_B^T \times (X_B \times \mathbf{w} - Y_B)$$

- We compute

$$\mathbf{w} = \mathbf{w} - \frac{1}{|B|} \alpha X_B^T \times (\phi(X_B \times \mathbf{w}) - Y_B)$$

- ϕ is very expensive to compute with just $+$, $*$ operations
- We build a piece-wise polynomial protocol for $\phi'(\cdot) \approx \phi(\cdot)$

$$\phi'(z) = \max\left(0, \min\left(z + \frac{1}{2}, 1\right)\right)$$



Logistic Regression

- Instead of computing

$$\mathbf{w} = \mathbf{w} - \frac{1}{|B|} \alpha X_B^T \times (X_B \times \mathbf{w} - Y_B)$$



- We compute

$$\mathbf{w} = \mathbf{w} - \frac{1}{|B|} \alpha X_B^T \times (\phi(X_B \times \mathbf{w}) - Y_B)$$

- ϕ is very expensive to compute with just $+$, $*$ operations
- We build a piece-wise polynomial protocol for $\phi'(\cdot) \approx \phi(\cdot)$

$$\phi'(z) = \max\left(0, \min\left(z + \frac{1}{2}, 1\right)\right)$$

or higher degree approximation




Neural Network


- Generalization of logistic regression
 - Each node is a regression problem
 - Replace logistic function with ReLU.
 - Easy to implement with piecewise polynomial
- Output nodes use soft-max:

$$\sigma_i(\vec{x}) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

- For inference/prediction, replace with arg-max.
- For training,
 - Approximate e^x
 - Approximate $1/x$ or use one garbled circuit



ReLU: $f(x) = \max(0, x)$



Performance – Inference

The models are for the MNIST dataset with $D = 784$ features.



Model	Protocol	Batch Size	Running Time (ms)		Comm. (MB)
			Online	Total	
Linear 93%	This	1	0.1	3.8	0.002
		100	0.3	4.1	0.008
	SecureML [47]	1	0.2	2.6	1.6
		100	0.3	54.2	160
Logistic* 98%	This	1	0.2	4.0	0.005
		100	6.0	9.1	0.26
	SecureML [47]	1	0.7	3.8	1.6
		100	4.0	56.2	161
NN 97%	This	1	3	8	0.5
	SecureML [47]	1	193	4823	120.5
CNN 99%	This*	1	6	10	5.2
	Chameleon [50]	1	1360	2700	12.9
	MiniONN [43]	1	3580	9329	657.5

Performance – Logistic Regression Training

- Measures iterations / second, **larger = better**
- Dimension = # of features
- Batch Size B = # examples used at each iteration
- Total running time: Up to **700 times faster**
- Communication: Up to **600 times less**

Dimension	Protocol	Batch Size B							
		Online				Online + Offline			
		128	256	512	1024	128	256	512	1024
10	This	2251	2053	1666	1245	2116	1892	1441	1031
	[47]	188	101	41	25	37	20	8.6	4.4
100	This	1867	1375	798	375	1744	1276	727	345
	[47]	183	93	46	24	3.6	1.9	1.1	0.6
1000	This	349	184	95	42	328	177	93	41
	[47]	105	51	24	13.5	0.43	0.24	0.12	0.06

Summary – ABY³

- New Protocols:
 - Efficient support for fixed-point arithmetics
 - Improved matrix multiplication
 - Efficient piece-wise polynomial evaluation
 - Conversions between **A**rithmetic, **B**oolean, and **Y**ao secret shares
- Prototype machine learning implementation:
 - Linear Regression – 12000 iterations / second
 - Logistic Regression – 2000 iterations / second
 - Neural Networks – 10 millisecond inference

The End, Questions?


Peter Rindal
Payman Mohassel

VISA
Research

Truncation on 2-out-of-2 Secret Shares

[MohasselZhang17]

Party 1



Assumption: no overflowing



Operation: truncate top/bottom bits




-1 error if [λ bits (hatched)] + [m bits (hatched)] overflows

$$\Pr[\text{large error}] = 2^{-\lambda}$$




= Probability of -1 error removing a mod 2^k wrap around
 $\leq 2^{-(\text{number of leading zeros})} = 2^{-\lambda}$

Party 2



Truncation on 2-out-of-3 Secret Shares

Party 1



Assumption: no overflowing



Operation: truncate top/bottom bits




-/ error if $\lambda + \lambda$ overflows

$$\Pr[\text{large error}] = 2^{-\lambda}$$





~~= Probability of 1 error removing a mod 2^k wrap around~~
 ~~$\leq 2^{-(\text{number of leading zeros})} = 2^{-\lambda}$~~

Party 2



Party 3



Truncation/Multiplication on 2-out-of-3 Secret Shares

- Multiplication $z = x * y$:

- Preprocess

- $\llbracket r' \rrbracket = rand()$

- $\llbracket r \rrbracket = \llbracket r' \rrbracket / d$

- Let $\llbracket z' \rrbracket = \llbracket x \rrbracket \llbracket y \rrbracket$ (integer multiplication)

$$z'_1 = x_1 y_1 + x_1 y_3 + x_3 y_1$$

$$z'_2 = x_1 y_2 + x_2 y_2 + x_2 y_2$$

$$z'_3 = x_2 y_3 + x_3 y_2 + x_3 y_3$$

- Compute

$$\llbracket z \rrbracket = \frac{reveal(\llbracket z' \rrbracket - \llbracket r' \rrbracket)}{d} + \llbracket r \rrbracket$$

VISA

$$\begin{matrix} x_1, y_1, z'_1 \\ x_3, y_3, z'_3 \end{matrix} - r'_1$$

$$z' - r'$$

$$\begin{matrix} x_1, y_1, \\ x_2, y_2, z'_2 \end{matrix} - r'_2$$

$$\begin{matrix} x_2, y_2, \\ x_3, y_3, z'_3 \end{matrix} - r'_3$$

Bank of America

amazon

Truncation/Multiplication on 2-out-of-3 Secret Shares

- Multiplication $z = x * y$:

- Preprocess

- $\llbracket r' \rrbracket = rand()$

- $\llbracket r \rrbracket = \llbracket r' \rrbracket / d$

- Let $\llbracket z' \rrbracket = \llbracket x \rrbracket \llbracket y \rrbracket$ (integer multiplication)

$$z'_1 = x_1 y_1 + x_1 y_3 + x_3 y_1$$

$$z'_2 = x_1 y_2 + x_2 y_2 + x_2 y_2$$

$$z'_3 = x_2 y_3 + x_3 y_2 + x_3 y_3$$

- Compute

$$\llbracket z \rrbracket = \frac{z' - r'}{d} + \frac{\llbracket r' \rrbracket}{d}$$

VISA

x_1, y_1, z'_1
 $x_3, y_3,$

$$z' - r'$$

$x_1, y_1,$
 x_2, y_2, z'_2

$x_2, y_2,$
 x_3, y_3, z'_3

Bank of America

amazon

Performance – Linear Regression Training

- Measures iterations / second, **larger = better**
- Dimension = # of features
- Batch Size B = # examples used at each iteration

Dimension	Protocol	Batch Size B							
		Online Throughput				Online + Offline Throughput			
		128	256	512	1024	128	256	512	1024
10	This	11764	10060	7153	5042	11574	9803	6896	4125
	[47]	7889	7206	4350	4263	47	25	11	5.4
100	This	5171	2738	993	447	5089	2744	1091	470
	[47]	2612	755	325	281	3.7	2.0	1.1	0.6
1000	This	406	208	104	46	377	200	100	46
	[47]	131	96	45	27	0.44	0.24	0.12	0.06