

MARbled Circuits: Mixing Arithmetic and Boolean Circuits with Active Security[★]

Dragos Rotaru^{1,2} and Tim Wood^{1,2}

¹ University of Bristol, Bristol, UK

² imec-COSIC KU Leuven, Leuven, Belgium

`dragos.rotaru@esat.kuleuven.be, t.wood@kuleuven.be`

Abstract. Most modern actively-secure multiparty computation (MPC) protocols involve generating random data that is secret-shared and authenticated, and using it to evaluate arithmetic or Boolean circuits in different ways. In this work we present a generic method for converting authenticated secret-shared data between different fields, and show how to use it to evaluate so-called “mixed” circuits with active security and in the full-threshold setting. A mixed circuit is one in which parties switch between different subprotocols dynamically as computation proceeds, the idea being that some protocols are more efficient for evaluating arithmetic circuits, and others for Boolean circuits.

One use case of our switching mechanism is for converting between secret-sharing-based MPC and garbled circuits (GCs). The former is more suited to the evaluation of arithmetic circuits and can easily be used to emulate arithmetic over the integers, whereas the latter is better for Boolean circuits and has constant round complexity. Much work already exists in the two-party semi-honest setting, but the n -party dishonest majority case was hitherto neglected.

We call the actively-secure mixed arithmetic/Boolean circuit a marbled circuit³. Our implementation showed that mixing protocols in this way allows us to evaluate a linear Support Vector Machine with 400 times fewer AND gates than a solution using GC alone albeit with twice the preprocessing required using only SPDZ (Damgård et al., CRYPTO ’12), and thus our solution offers a tradeoff between online and preprocessing complexity. When evaluating over a WAN network, our online phase is 10 times faster than the plain SPDZ protocol.

1 Introduction

One of the major modern uses of cryptography is for **mutually-distrustful** parties to compute a function on their combined secret inputs so that all parties learn the output and no party learns anything more about other parties’ inputs than what can be deduced from their own input and the output alone. This is known as secure multiparty computation (MPC) and has recently been shown to be very efficient for evaluating general Boolean [NNOB12, DZ13] and arithmetic [DPSZ12, DKL⁺13, KOS16, KPR18] circuits.

Many real-world use cases of computing on private data involve some form of statistical analysis, requiring evaluation of arithmetic formulae. Perhaps the most common method of computing arithmetic circuits on private data involves *secret-sharing* (SS), in which secret inputs are split up into several pieces and distributed amongst a set of parties, which then

[★] This work has been supported in part by ERC Advanced Grant ERC-2015-AdG-IMPACT, by the Defense Advanced Research Projects Agency (DARPA) and Space and Naval Warfare Systems Center, Pacific (SSC Pacific) under contract No. N66001-15-C-4070, and by the FWO under an Odysseus project GOH9718N.

³ See paper marbling.

perform computations on these shares (sometimes requiring communication), and recombine them at the end for the result. However, MPC over a finite field or a ring is used to emulate arithmetic over the integers, and consequently, operations such as comparisons between secrets (i.e. $<$, $>$, $=$), which we refer to generally as “bit-wise” operations, are an important feature of MPC protocols: one of the shortcomings of MPC based on secret-sharing is that these natural but more complicated procedures require special preprocessing and several rounds of communication.

One way to mitigate these costs would be to use *garbled circuits* (GCs) instead of secret-sharing for circuits involving lots of non-linear operations, since this method has lower round complexity than SS-based MPC solutions (in fact, they can be done in constant rounds). Recent work has shown that multiparty Boolean circuit garbling with active security in the dishonest majority setting can be made very efficient [WRK17, HSS17, KY18]. However, performing general arithmetic computations in Boolean circuits can be expensive since the arithmetic operations must be accompanied by reduction modulo a prime inside the circuit. Moreover, efficient constructions of multiparty constant-round protocols for *arithmetic* circuits remain elusive. Indeed, the best-known optimisations for arithmetic circuits such as using a primordial modulus [BMR16] are expensive even for passive security in the two-party setting. The only work of which the authors are aware in the multiparty setting is the passively-secure honest-majority work by Ben-Efraim [Ben18].

So-called *mixed protocols* are those in which parties switch between secret-sharing (SS) and a garbled circuit (GC) mid-way through a computation, thus enjoying the efficiency of the basic addition and multiplication operations in any field using the former and the low-round complexity of GCs for more complex subroutines using the latter. One can think of mixed protocols as allowing parties to choose the most efficient field in which to evaluate different parts of a circuit.

There has been a lot of work on developing mixed protocols in the two-party passive security setting, for example [HKS⁺10, KSS13, KSS14, BDK⁺18]. One such work was the protocol of Demmler et al. [DSZ15] known as ABY, that gave a method for converting between arithmetic, Boolean, and Yao sharings. For small subcircuits, converting arithmetic shares to Boolean shares (of the bit decomposition) of the same secret – i.e. without any garbling – was shown to give efficiency gains over performing the same circuits in with arithmetic shares; for large subcircuits, using garbling allows reducing online costs. Mohassel and Rindal [MR18] constructed a three-party protocol known as ABY3 for mixing these three types of sharing in the malicious setting assuming at most one corruption.

For mixed protocols to be efficient, clearly the cost of switching between secret-sharing and garbling, performing the operation, and switching back must be more efficient than the method that does not require switching, perhaps achieved by relegating some computation to the offline phase.

1.1 Our Contribution

The main challenge for active security is that it is essential to retain authentication of secrets through the conversion. In this work, we give a simple actively-secure procedure for

transforming data that is secret-shared and authenticated in different ways. **The motivation is to** allow mixed protocols in the dishonest majority setting with active security, with only black-box use of the linear secret sharing scheme (LSSS) and GC subprotocols. The idea behind specifically designing a transformation procedure instead of a whole MPC protocol is that any circuit that makes considerable use of *both* arithmetic operations *and* bit-wise computations is likely to be most efficient when using state-of-the-art SS-based MPC and circuit garbling protocols, assuming the transformation procedure is cheap enough. Our implementation shows that this is achievable with concrete efficiency, and that there is some tradeoff between preprocessing costs and circuit evaluation costs. In the following discussion, we will focus on specific goal of switching between secret-sharing schemes and GCs.

Let \mathbb{F}_q denote the finite field of order q . One of the key observations that allows our generic transformation to be realised is that for many recent protocols in both the SS-based and GC-based MPC paradigms, the starting point is essentially always to create a black-box actively-secure SS-based MPC functionality and to use it either directly to evaluate arithmetic circuits if the field is \mathbb{F}_p , or to generate garbled circuits with active security if the field is \mathbb{F}_2 . At the highest level, the idea of our protocol is to allow data embedded in \mathbb{F}_p to be efficiently transformed into data embedded in \mathbb{F}_2 *with authentication*. It is then easy to show that garbled circuits can be used to evaluate on authenticated elements of \mathbb{F}_p , as will be demonstrated in this work.

The most obvious way of providing inputs that are secret-shared into a garbled circuit is relatively straightforward: for a given secret, parties can simply input each bit in the bit-decomposition of their share into the GC, and the sum mod p can be computed inside the circuit. The primary technical challenge for a conversion procedure with active security is to maintain *authentication* through the transition from secret-shared inputs and secret inputs inside the GC, and *vice versa*. The naïve way of obtaining authentication is for parties to bit-decompose the shares of the data that provides authentication validating inside the GC: for example, if information-theoretic MACs are used on secrets, parties input their shares of the MACs and the MAC key(s). Final circuit outputs can be sets of bits that represent the bit-decomposition of in \mathbb{F}_p , and can be privately opened to different parties. This method requires garbling several additions and multiplications inside the circuit to check the MACs and would require $O(n \cdot \kappa \cdot \log |\mathbb{F}|)$ bits per party to be sent to switch inputs in the online phase, where n is the number of parties, κ is the computational security parameter, and \mathbb{F} is the MPC field, since each party needs to broadcast a GC key for each bit of the input. The advantage of this solution, despite these challenges, is that it requires no additional preprocessing, nor adaptations to the garbling procedure.

Contrasting this approach, our solution makes use of special preprocessing to speed up the conversion. This results in reducing the circuit size by approximately 100,000 AND gates per conversion for a field with a 128-bit prime modulus (assuming Montgomery multiplication is used). In this work we show how to convert between secret-shared data in \mathbb{F}_p , where p is a large prime and is the MPC modulus, and GCs in \mathbb{F}_{2^k} through the use of “double-shared” **authenticated bits** which we dub *daBits*, following the nomenclature set out by [NNOB12]. These doubly-shared secrets are values in $\{0, 1\}$ shared and authenticated both in \mathbb{F}_p and

\mathbb{F}_{2^k} , where by 0 and 1 we mean the additive and multiplicative identity, respectively, in each field. **In brief**, the conversion of a SS input x into a GC involves constructing a random secret r in \mathbb{F}_p using daBits, opening $x - r$ in MPC, bit decomposing this public value (requiring no communication) and using these as signal bits for the GC, and then in the circuit adding r and computing this modulo p , which is possible since the bit decomposition of r is doubly-shared. This keeps the authentication check outside of the circuit instead requiring that the check happens correctly on the opened value $x - r$. Going the other way around, the output of the circuit is a set of public signal bits whose masking bits are chosen to be daBits. To get the output, parties XOR the public signal bits with the \mathbb{F}_p shares of the corresponding daBit masks, which can be done locally. These shares can then be used to reconstruct elements of \mathbb{F}_p (or remain as bits if desired).

The only use of doubly-shared masks is at the two boundaries (input and output) between a garbled circuit and secret-sharing; all secrets used in evaluating arithmetic circuits (i.e. using standard SS-based MPC) are authenticated shares in \mathbb{F}_p only; all wire masks “inside” the circuit (that is, for all wires that are *not* input or output wires) are authenticated shares of bits in \mathbb{F}_{2^k} only. The *online* communication cost of our solution is that of each party broadcasting a single field element and then broadcasting $\log |\mathbb{F}|$ key shares per input, for a circuit of any depth. Thus the cost is $O(\kappa \log |\mathbb{F}|)$ per party, **per** field input to the circuit. The offline cost grows quadratically in n as generating daBits requires every party to communicate with every other party.

We emphasise that while we focus on allowing conversion between SS and GCs, the conversion is generic in the sense that it is merely a method of converting data embedded into one field into data embedded in another field, with authentication attached. For example, once the bits of $x - r$ are public and the bit decomposition of r is known in \mathbb{F}_{2^k} , parties can execute an SS-based MPC protocol on these bits directly, without going through garbled circuits. Thus our work is also compatible with converting classic SPDZ shares in \mathbb{F}_p with the recent protocol SPDZ2k of Cramer et al. [CDE+18].

We remark that several of the multiparty arithmetic garbling techniques of [Ben18] require the use of “multifield shared bits”, which precisely correspond to our daBits (albeit in an unauthenticated honest-majority setting), and consequently we suggest that our idea of generating daBits may lead to more efficient actively-secure multiparty garbling of arithmetic circuits.

Related work. Recent work [KSS13, KSS10] deals with conversion between homomorphic encryption (HE) and GC for two parties. In their two-party case, the conversion works by having P_1 encrypt a blinded version of its input $x - r$ using the public key of P_2 . Since the second party can decrypt the ciphertext, the share of x held by P_2 is $x - r$ and the share of P_1 is r . They also provide some techniques to convert an additive sharing in a ring to a GC sharing in the malicious case (P_2 only). Unfortunately their solution is insecure in the case of a malicious P_1 without adding some extra zero knowledge proofs. Moreover, it is unclear how to extend their work to convert from an n -party authenticated LSSS sharing of x to an n -party sharing inside a GC with a dishonest majority.

Active security beyond bounded inputs While essentially all of the basic actively-secure MPC protocols enable the evaluation of additions and multiplications, for more complicated non-linear functions the only solutions that exist are those that require additional assumptions on the input data. For example, comparison requires bit decomposition, which itself requires that all secrets be bounded by some constant. Since the bits of each input are directly inserted into the circuit, we can avoid this additional assumption. We refer the reader to [DFK⁺06] or the documentation for the SCALE/MAMBA project [AKO⁺18, §10 Advanced Protocols] for an overview of implementations of other functions in MPC.

2 Preliminaries

In this section we explain the basics of MPC as required to understand the sequel. In our protocol, one instance of MPC is used to perform the secret-sharing-based MPC over a prime field, and another instance is used to perform another form of MPC – typically, circuit garbling – over a large field of characteristic 2.

2.1 General

We denote the number of parties by n , and the set of indices of parties corrupted by the adversary by A . We write $[j]$ to denote the set $\{1, \dots, j\}$. We write \mathbb{F} to denote a field, and \mathbb{F}_q to denote the finite field of order q . The arithmetic circuit will be computed in the field \mathbb{F}_p where p is a large prime, and the keys and masks for the garbled circuit in \mathbb{F}_{2^k} . By $\log(\cdot)$ we always mean the base-2 logarithm, $\log_2(\cdot)$. We denote by sec and κ the statistical and computational security parameters, respectively. We say that a function $\nu : \mathbb{N} \rightarrow \mathbb{R}$ is *negligible* if for every polynomial $f \in \mathbb{Z}[X]$ there exists $N \in \mathbb{N}$ such that $|\nu(X)| < |1/f(X)|$ for all $X > N$. If an event X happens with probability $1 - \nu(\text{sec})$ where ν is a negligible function then we say that X happens with overwhelming probability in sec . We write $x \xleftarrow{\$} S$ to mean that x is sampled uniformly from the set S , and use $x \leftarrow y$ to mean that x is assigned the value y . We will sometimes denote by, for example, $(a - b)_j$ the j^{th} bit in the binary expansion of the integer $a - b$.

2.2 Security

UC Framework We prove our protocols secure in the universal composability (UC) framework of Canetti [Can01]. Protocols proved secure in this model are secure even when executed alongside arbitrarily many other protocols, concurrently, sequentially or both. We assume the reader’s familiarity with this framework. In Figure 1 we give a functionality $\mathcal{F}_{\text{Rand}}$ for obtaining unbiased random data that we need for our protocol. A protocol realising $\mathcal{F}_{\text{Rand}}$, and other UC functionalities, are given in Appendix B.

Functionality $\mathcal{F}_{\text{Rand}}$

Random subset On input $(\text{RSubset}, X, t)$ where X is a set satisfying $|X| \geq t$, sample $S \xleftarrow{\$} \{A \subseteq X : |A| = t\}$ and send S to all parties.

Random buckets On input $(\text{RBucket}, X, t)$ where X is a set and $t \in \mathbb{N}$ such that $|X|/t \in \mathbb{N}$, set $n \leftarrow |X|/t$ and then for each $i = 1, \dots, n$ do the following:

1. Sample $X_i \xleftarrow{\$} \{A \subseteq X : |A| = t\}$.
2. Set $X \leftarrow X \setminus X_i$.

Finally, send $(X_i)_{i=1}^n$ to all parties.

Fig. 1. Functionality $\mathcal{F}_{\text{Rand}}$

We assume an active, static adversary corrupting at most $n - 1$ out of n parties. Adversaries corrupting at most all parties but one are called “full-threshold”. An active adversary may deviate arbitrarily from the protocol description, and a static adversary is permitted to choose which parties to corrupt only at the beginning of the protocol, and cannot corrupt more parties later on. Our protocol allows corrupt parties to cause the protocol to abort before honest parties receive output, but if the adversary cheats then the honest parties will not accept an incorrect output. This is known as “security-with-abort” in the literature. While this work focuses on the full-threshold setting, since \mathcal{F}_{MPC} is used as a black box, the access structure of our protocol is solely dependent on the access structure admitted by the instantiation(s) of \mathcal{F}_{MPC} - for a complete definition of \mathcal{F}_{MPC} which is sometimes denoted as \mathcal{F}_{ABB} or $\mathcal{F}_{\text{AMPC}}$ in the literature check [BDOZ11, SW19].

Communication We assume point-to-point secure channels, and synchronous communication. Additionally, we assume a broadcast channel, which can be instantiated in the random oracle model over point-to-point secure channels, for example as described in [DPSZ12, App. A.3]. A round of communication is a period of time in which parties perform computation and then send and receive messages. Messages sent in a round cannot depend on messages received during the current round, but messages may depend on messages sent in previous rounds. So-called *constant-round* protocols require $O(1)$ rounds for the entire protocol.

2.3 Garbled Circuits

Essentially all modern dishonest-majority multiparty Boolean circuit garbling protocols, for example [HSS17, WRK17, HOSS18], use some form of secret-sharing based MPC to generate a multiparty garbled circuit, thus following the basic idea of SPDZ-BMR [LPSY15] of using MPC to ensure correctness of the garbling in the classic multiparty garbling protocol by Beaver et al. [BMR90]. The specific garbling protocol is not important for this work: our solution makes use of the garbling subprotocols in a black-box way. For this reason, the explanation of MPC, with authentication for active security, is given below; the explanation of the garbling subprotocols is left to Appendix A. The important parts of the garbling evaluation are providing inputs from \mathbb{F}_p into a generic multiparty garbled circuit, and extracting \mathbb{F}_p outputs from it, which is dealt with in Section 4.2.

2.4 MPC

Our protocol **makes use of** MPC **as** a black box, for which the generic functionality is outlined in Figure 3 as part of the functionality $\mathcal{F}_{\text{Prep}}$. The functionality \mathcal{F}_{MPC} over a field \mathbb{F} is realised using protocols with statistical security sec if $|\mathbb{F}| = \Omega(2^{\text{sec}})$ and computational security κ depending on the computational primitives being used. We will describe MPC as executed in the SPDZ family of protocols [DPSZ12, DKL+13, KOS16, KPR18, CDE+18]. These protocols are in the preprocessing model in which circuit evaluation is split into a *preprocessing* (or *offline*) phase in which input-independent data is generated, that is then “**used up**” in an *online* phase which uses the actual circuit inputs. The reason for doing this is that the preprocessed data is expensive to generate, but the online evaluation is cheap as a result.

Note that MPC as used in garbling often offers “less” than full circuit evaluation as described here: for example, [HSS17] defined a protocol $\Pi_{\text{Bit} \times \text{String}}$ that allows bits to be multiplied by bitstrings but is not concerned with multiplication of general field elements in \mathbb{F}_{2^k} . However, what follows is enough to understand the main techniques required for SS-based multiparty garbling.

Secret-sharing A secret $x \in \mathbb{F}$ is said to be *additively shared*, denoted by $\llbracket x \rrbracket$, if a dealer samples a set $\{x^i\}_{i=1}^{n-1} \xleftarrow{\$} \mathbb{F}$, fixes $x^n \leftarrow x - \sum_{i=1}^{n-1} x^i$, and for all $i \in [n]$ sends x^i to party P_i . Any set of $n-1$ shares is indistinguishable from a set of $n-1$ uniformly-sampled shares, and the sum of all n shares is the secret x . This secret-sharing is linear: the sum of corresponding shares of two secrets is a sharing of the sum of the two secrets, so no communication is required for linear functions. A secret x shared in this way is denoted by $\llbracket x \rrbracket = (x^i)_{i=1}^n$.

MPC protocols based on secret-sharing involve secret-sharing all the secret inputs, performing computations on the shares, and recombining at the end to obtain the final result.

Addition of secrets Addition of secrets is denoted as follows:

$$\llbracket a \rrbracket + \llbracket b \rrbracket \leftarrow \llbracket a + b \rrbracket = (a^i + b^i)_{i=1}^n$$

Thus any linear function on secret-shared data can be evaluated without communication.

Multiplication of secrets Using a technique due to Beaver [Bea92], multiplication of secrets in the online phase can be computed as a *linear* operation if the parties have access to so-called *Beaver triples* – triples of secret-shared values $(\llbracket a \rrbracket, \llbracket b \rrbracket, \llbracket a \cdot b \rrbracket)$, where a and b are uniformly random and unknown to the parties. To multiply $\llbracket x \rrbracket$ and $\llbracket y \rrbracket$, the parties compute $\llbracket x - a \rrbracket \leftarrow \llbracket x \rrbracket - \llbracket a \rrbracket$ and $\llbracket y - b \rrbracket \leftarrow \llbracket y \rrbracket - \llbracket b \rrbracket$ locally and open them (i.e. they broadcast their shares of $\llbracket x - a \rrbracket$ and $\llbracket y - b \rrbracket$ so all parties learn $x - a$ and $y - b$), and then compute the product as

$$\llbracket x \cdot y \rrbracket \leftarrow \llbracket a \cdot b \rrbracket + (x - a) \cdot \llbracket b \rrbracket + (y - b) \cdot \llbracket a \rrbracket + (x - a) \cdot (y - b).$$

Since a and b are unknown to any party, $x - a$ and $y - b$ reveal nothing about x and y . We refer to protocols that generate Beaver triples as *SPDZ-like*. **The main cost of SPDZ-like protocols comes from generating these Beaver triples.** The two main ways of doing this

are either using somewhat homomorphic encryption (SHE) [DPSZ12, DKL⁺13, KPR18] or oblivious transfer (OT) [KOS16, CDE⁺18].

Authentication Since both addition and multiplication involve only linear operations, computations on shares in the online phase only need to be protected against *additive* errors – that is, a corrupt party P_i changing its share from x_i to $x_i + \varepsilon$. To protect against such errors, linear information-theoretic MACs are used. Since these MACs are linear, parties can maintain MACs on every secret throughout the whole circuit evaluation; then they can run an amortised check of their correctness once at the end of the protocol execution. Though we are not concerned with the specifics of authentication as used in SS-based GC or general MPC in this work, it is helpful to understand how secrets can be authenticated in different ways, and so a brief description of two prevalent forms follows.

SPDZ-style MACs. A secret $a \in \mathbb{F}$ is shared amongst the parties by additively sharing the secret a in \mathbb{F} along with a linear MAC $\gamma(a)$ defined as $\gamma(a) \leftarrow \alpha \cdot a$, where $\alpha \in \mathbb{F}$ is a global MAC key, which is also additively shared. By “global” we mean that every MAC in the protocol uses this MAC key, rather than each party holding their own key and authenticating every share held by every other party. Note that the parties can trivially obtain a MAC on a public value a by computing $\llbracket \gamma(a) \rrbracket \leftarrow (\alpha^i \cdot a)_{i=1}^n$.

Now if p is $O(2^{\text{sec}})$ then to introduce an error ε on the sharing requires modifying the corresponding MAC by $\varepsilon \cdot \alpha$ – i.e. the adversary must guess the MAC key. We refer the reader to [DKL⁺13] for details on the MAC checking procedure.

BDOZ-style MACs. A different type of MAC, sometimes called *pairwise*, used by Bendlin et al. [BDOZ11]. A bit $c \in \mathbb{F}_2$ is shared as $c = \bigoplus_{i=1}^n c^i$ where for each share $c^i \in \mathbb{F}_2$ (held by P_i), for each $j \neq i$, parties P_i and P_j hold a MAC as follows: P_i holds $M_j^i[c] \in \mathbb{F}_{2^k}$ and P_j holds $K_i^j[c] \in \mathbb{F}_{2^k}$ such that $K_i^j[c] \oplus M_j^i[c] = \Delta^j \cdot c^i$, where Δ^j is the MAC key held by P_j and c^i a random bit held by P_i . This MAC scheme is used in the garbling protocols due to Hazay et al. and Wang et al. [HSS17, WRK17].

Sharing Notation Below we give the precise meaning of the notation $\llbracket a \rrbracket_p$ and $\llbracket c \rrbracket_{2^k}$. One can think of using SPDZ-style MACs for \mathbb{F}_p and BDOZ-style MACs for \mathbb{F}_{2^k} , as described below, but the protocols for conversion are oblivious to the precise method of authentication so these choices are essentially arbitrary. The third type of sharing is the notation that will be used for our special preprocessing called daBits.

$$\text{LSSS Sharing} \quad \llbracket a \rrbracket_p = (a^i, \gamma_p(a)^i, \alpha^i)_{i=1}^n$$

$$\text{GC Sharing} \quad \llbracket c \rrbracket_{2^k} = (c^i, (K_i^j[c])_{j \neq i}, (M_j^i[c])_{j \neq i}, \Delta^i)_{i=1}^n.$$

$$\text{Sharing in both} \quad \llbracket b \rrbracket_{p,2^k} = (\llbracket b \rrbracket_p, \llbracket b \rrbracket_{2^k}) \text{ where } b \in \{0, 1\}.$$

The sharing $\llbracket b \rrbracket_{p,2^k}$ is considered correct if the bit is the same in both fields (either 0 or 1). Creating these bits while guaranteeing active security is one of the contributions of this work.

Conditions on the secret-sharing field Let $l = \lfloor \log p \rfloor$. Throughout, we assume the MPC is over \mathbb{F}_p where p is some large prime, but we require that one must be able to generate uniformly random field elements by sampling bits uniformly at random $\{\llbracket r_j \rrbracket_p\}_{j=0}^{l-1}$ and summing them to get $\llbracket r \rrbracket_p \leftarrow \sum_{j=0}^{l-1} 2^j \cdot \llbracket r_j \rrbracket_p$. For this to hold in \mathbb{F}_p , we require that $\frac{p-2^l}{p} = O(2^{-\text{sec}})$. Roughly speaking this says that p is slightly larger than a power of 2. (By symmetry of this argument we can require that p be close to a power of 2.) It follows from Lemma 2.1 that the statistical distance between the uniform distribution over \mathbb{F}_p and the same over $\{0, 1\}^l$ is negligible.

Lemma 2.1. *Let $l = \lfloor \log p \rfloor$, let P be the probability mass function for the uniform distribution \mathcal{P} over $[0, p) \cap \mathbb{Z}$ and let Q be the probability mass function for the uniform distribution \mathcal{Q} over $[0, 2^l) \cap \mathbb{Z}$. Then the statistical distance between distributions is negligible in the security parameter if $\frac{p-2^l}{p} = O(2^{-\text{sec}})$.*

Proof. By definition of statistical distance,

$$\begin{aligned} \Delta(\mathcal{P}, \mathcal{Q}) &= \frac{1}{2} \cdot \sum_{x=0}^{p-1} |P(x) - Q(x)| = \frac{1}{2} \cdot \sum_{x=0}^{2^l-1} \left| \frac{1}{p} - \frac{1}{2^l} \right| + \frac{1}{2} \cdot \sum_{x=2^l}^{p-1} \left| \frac{1}{p} - 0 \right| \\ &= \frac{1}{2} \cdot 2^l \cdot \frac{p-2^l}{p \cdot 2^l} + \frac{1}{2} \cdot (p-2^l) \cdot \frac{1}{p} = \frac{p-2^l}{p} = O(2^{-\text{sec}}). \end{aligned}$$

Notation The functionalities sometimes refer to identifiers for variables such as `id`, where the value inside is `Val[id]`. We will use $\llbracket x \rrbracket$ to denote the variable identifier `id` for the value x , so $\text{Val}[\llbracket x \rrbracket] = x$, and saying that the parties have $\llbracket x \rrbracket$ does not imply they know the secret x . This intentional collision of notation for authenticated secrets is to demonstrate that the realisation of the dictionary `Val` occurs via the secret sharing with MACs. To save on overloaded notation, we will occasionally write $\llbracket z \rrbracket \leftarrow \llbracket x \rrbracket + \llbracket y \rrbracket$ to mean that parties send the command $(\text{Add}, \llbracket x \rrbracket, \llbracket y \rrbracket, \llbracket z \rrbracket)$ to some functionality, and similarly for multiplication. Where not explicitly specified, new identifiers are taken from a counter which the parties initialise to 0 at the start of the protocol.

Note on XOR In our context, we will require heavy use of the (generalised) XOR operation. This can be defined in any field as the function

$$\begin{aligned} f : \mathbb{F}_p \times \mathbb{F}_p &\rightarrow \mathbb{F}_p \\ (x, y) &\mapsto x + y - 2 \cdot x \cdot y, \end{aligned} \tag{1}$$

which coincides with the usual XOR function for **fields of characteristic 2**. In SS-based MPC, addition requires no communication, so computing XOR in \mathbb{F}_{2^k} is for free; the cost in \mathbb{F}_p ($p > 2$) is one multiplication, which requires preprocessed data and some communication. This operation turns out to be the main cost associated with our offline phase (see Table 2).




Fig. 2. Share conversions for dishonest majority protocols. Dashed lines use our daBits as an inner subroutine.

3 Generality of daBits

In Diagram 2 we show how our daBit generation bridges different MPC protocols for dishonest majority. Our inspiration is drawn from Keller and Yanai [KY18] which can convert between SPDZ-BMR and SPDZ over \mathbb{F}_2^k by setting the global difference used in the free-XOR as the global MAC-key in $\text{SPDZ}[\mathbb{F}_2^k]$. Their main idea is to sample a secret random bit authenticated in \mathbb{F}_2^k and use that random bit to do a share conversion. This lends nicely because the authentication key has the same representation in both engines whereas we need more involved techniques (eg: use cut and choose) to generate such a preprocessed authenticated random bit between $\text{SPDZ}[\mathbb{F}_p]$ and SPDZ-BMR (or $\text{BMR}[\mathbb{F}_2^k]$).

A note on rings and fields. Our protocol uses actively-secure MPC as black box, so there is no reason the MPC cannot take place over any ring $\mathbb{Z}/m\mathbb{Z}$ where m is possibly composite, as long as m is (close to) a power of 2. The security of our procedure for generating daBits can tolerate zero-divisors in the ring, so computation may, for example, take place over the ring $\mathbb{Z}/2^l\mathbb{Z}$ for any l , for which actively-secure \mathcal{F}_{MPC} can be realised using [CDE⁺18].

BMR and TinyOT. This is done by converting the pairwise MAC to a global one and it is explained in several papers [HSS17, WRK17]. Going from a global MAC to a pairwise one is slightly more difficult but could be achieved using daBits or a similar consistency check by Damgård et al. [DEF⁺19] to go from a bit share in $\text{SPDZ}[\mathbb{Z}_{2^k}]$ to a TinyOT sharing.

TinyOT and SPDZ2k. Recently Damgård et al. [DEF⁺19] introduced a method of switching back and forth between $\text{SPDZ}[\mathbb{Z}_{2^k}]$ to TinyOT. They use a lightweight batch-check to ensure input consistency in both engines. Although they show how to switch a random bit $\llbracket b \rrbracket_{2^k} \in \mathbb{Z}_{2^k}$ to $\llbracket b \rrbracket_2 \in \mathbb{F}_2^k$ their subroutines can be used to convert a full input $\llbracket x \rrbracket_{2^k} \in \mathbb{Z}_{2^k}$ by bit-decomposing it and then translate each bit into the TinyOT family of protocols.

SPDZ2k and SPDZ. One can use daBits to convert from a $\text{SPDZ}[\mathbb{F}_p]$ share to a $\text{SPDZ}[\mathbb{Z}_{2^k}]$ share. The high level idea of converting between a field and a ring is to generate the same correlated randomness $\llbracket r \rrbracket_p \in \mathbb{F}_p$ and $\llbracket r \rrbracket_{2^k} \in \mathbb{Z}_{2^k}$ by bit-composing the daBits. Then the conversion from $\llbracket x \rrbracket_p$ to $\llbracket x \rrbracket_{2^k}$ becomes trivial: parties open $\llbracket x \rrbracket_p - \llbracket r \rrbracket_p$, assign this to a public y then perform the reduction modulo p in $\text{SPDZ}[\mathbb{Z}_{2^k}]$ using the public constant

y i.e. $\llbracket x \rrbracket_{2^k} \leftarrow (y + \llbracket r \rrbracket_{2^k}) \bmod p$. This procedure can be adapted to allow conversions from $\text{SPDZ}[\mathbb{Z}_{2^k}]$ to $\text{SPDZ}[\mathbb{F}_p]$: parties open $\llbracket x \rrbracket_{2^k} - \llbracket r \rrbracket_{2^k}$ in \mathbb{Z}_{2^k} then add the randomness back in \mathbb{F}_p and truncate the result modulo \mathbb{Z}_{2^k} . A similar idea can be applied to convert from $\text{SPDZ}[\mathbb{F}_p]$ to $\text{SPDZ}[\mathbb{F}_2^k]$ with the exception that in the last step parties now truncate their shares modulo \mathbb{F}_2^k .

SPDZ and TinyOT family. This conversion can be done again with daBits and works in the same way we describe it in our paper for $\text{SPDZ}[\mathbb{F}_p]$ and $\text{BMR}[\mathbb{F}_2^k]$. Recently Aly et al. [AOR⁺19] improve and fully integrate the conversion between SPDZ and WRK/HSS garbling into SCALE-MAMBA. Their protocol improvements come from a slightly modified check of Damgård et al. [DEF⁺19] with a twist in how parties extract the least significant bit of a SPDZ share by tweaking their shares locally in the two-party case. They achieve an amortized cost of just one \mathbb{F}_p triple per daBit.

SPDZ and BMR. This is the focus of our paper: parties generate multiple daBits $\llbracket b \rrbracket_{p,2^k}$ to form correlated randomness shared in both fields: $\llbracket r \rrbracket_{p,2^k} \leftarrow 2^i \llbracket b_i \rrbracket_{p,2^k}$. Afterwards they open their masked $\text{SPDZ}[\mathbb{F}_p]$ shares with $\llbracket r \rrbracket_p$ and remove the randomness inside BMR using $\llbracket r \rrbracket_{2^k}$. In the next section we describe how to produce this preprocessed material called daBits between $\text{SPDZ}[\mathbb{F}_p]$ and $\text{SPDZ-BMR}[\mathbb{F}_2^k]$ which is BMR and where the secret bit multiplications for the garbling are done using MASCOT [KOS16, KY18]. We now describe the protocol in producing these daBits.

4 Protocol

The idea behind creating a mixed protocol is to use one instance of \mathcal{F}_{MPC} over \mathbb{F}_p to perform addition and multiplication in the field, and one instance of \mathcal{F}_{MPC} over \mathbb{F}_{2^k} is used to perform the garbling, and to allow conversion between the two. Thus the goal is the functionality $\mathcal{F}_{\text{Prep}}$, which is given in Figures 3 and 4.

Note that since the keys for the PRF live in the field \mathbb{F}_{2^k} in the garbling protocol, the instance of \mathcal{F}_{MPC} must be over a field with $k = O(\kappa)$ for computational security. Indeed, we emphasise that in our protocol k is not directly related to $\log p$. Once the garbling is completed, the full MPC engine in \mathbb{F}_{2^k} is no longer required: the parties only maintain the \mathbb{F}_p instance of \mathcal{F}_{MPC} and retain the garbled circuits in memory, and will additionally need to make sure they can still perform the procedure **Check** in \mathcal{F}_{MPC} on values opened in the evaluation of the GC.

Recall from the introduction that the high-level idea of our protocol is to open a secret-shared value, locally bit-decompose it, and use these bits as input bits to the garbled circuit. Once the parties have these, they reveal corresponding secret pseudorandom function keys for circuit evaluation, and the rest of the protocol (including retrieving outputs in secret-shared form) is local.

Functionality $\mathcal{F}_{\text{Prep}}$

Instances of \mathcal{F}_{MPC}

Independent copies of \mathcal{F}_{MPC} are identified via session identifiers sid ; $\mathcal{F}_{\text{Prep}}$ maintains one dictionary Val_{sid} for each instance. Entries cannot be changed, for simplicity. If a party provides an input with an sid which has not been initialised, output *reject* to all parties and awaits another message.

Initialise On input $(\text{Initialise}, \mathbb{F}, \text{sid})$ from all parties, if sid is a new session identifier then initialise a database of secrets Val_{sid} indexed by a set $\text{Val}_{\text{sid}}.\text{Keys}$ and store the field as $\text{Val}_{\text{sid}}.\text{Field} \leftarrow \mathbb{F}$. Set the internal flag $\text{Abort}_{\text{sid}}$ to false.

Input On input $(\text{Input}, i, \text{id}, x, \text{sid})$ from P_i and $(\text{Input}, i, \text{id}, \perp, \text{sid})$ from all other parties, if $\text{id} \notin \text{Val}_{\text{sid}}.\text{Keys}$ then insert it and set $\text{Val}_{\text{sid}}[\text{id}] \leftarrow x$. Then call the procedure **Wait**.

Add On input $(\text{Add}, \text{id}_x, \text{id}_y, \text{id}, \text{sid})$, if $\text{id}_x, \text{id}_y \in \text{Val}_{\text{sid}}.\text{Keys}$ then set $\text{Val}_{\text{sid}}[\text{id}] \leftarrow \text{Val}_{\text{sid}}[\text{id}_x] + \text{Val}_{\text{sid}}[\text{id}_y]$.

Multiply On input $(\text{Multiply}, \text{id}_x, \text{id}_y, \text{id}, \text{sid})$, if $\text{id}_x, \text{id}_y \in \text{Val}_{\text{sid}}.\text{Keys}$ then set $\text{Val}_{\text{sid}}[\text{id}] \leftarrow \text{Val}_{\text{sid}}[\text{id}_x] \cdot \text{Val}_{\text{sid}}[\text{id}_y]$. Then call the procedure **Wait**.

Random element On input $(\text{RElt}, \text{id}, \text{sid})$, if $\text{id} \notin \text{Val}_{\text{sid}}.\text{Keys}$ then set $\text{Val}_{\text{sid}}[\text{id}] \xleftarrow{\$} \text{Val}_{\text{sid}}.\text{Field}$. Then call the procedure **Wait**.

Random bit On input $(\text{RBit}, \text{id}, \text{sid})$, if $\text{id} \notin \text{Val}_{\text{sid}}.\text{Keys}$ then set $\text{Val}_{\text{sid}}[\text{id}] \xleftarrow{\$} \{0, 1\}$. Then call the procedure **Wait**.

Open On input $(\text{Open}, i, \text{id}, \text{sid})$ from all parties, if $\text{id} \in \text{Val}_{\text{sid}}.\text{Keys}$,

- if $i = 0$ then send $\text{Val}_{\text{sid}}[\text{id}]$ to the adversary and run the procedure **Wait**. If the message was (OK, sid) , await an error ε from the adversary. Send $\text{Val}_{\text{sid}}[\text{id}] + \varepsilon$ to all honest parties and if $\varepsilon \neq 0$, set the internal flag $\text{Abort}_{\text{sid}}$ to true.
- if $i \in A$, then send $\text{Val}_{\text{sid}}[\text{id}]$ to the adversary and then run **Wait**.
- if $i \in [n] \setminus A$, then call the procedure **Wait**, and if not already halted then await an error ε from the adversary. Send $\text{Val}_{\text{sid}}[\text{id}] + \varepsilon$ to P_i and if $\varepsilon \neq 0$ then set the internal flag $\text{Abort}_{\text{sid}}$ to true.

Check On input $(\text{Check}, \text{sid})$ from all parties, run the procedure **Wait**. If not already halted and the internal flag $\text{Abort}_{\text{sid}}$ is set to true, then send the message $(\text{Abort}, \text{sid})$ to the adversary and honest parties and ignore all further messages to \mathcal{F}_{MPC} with this sid ; otherwise send the message (OK, sid) and continue.

Internal procedure:

Wait Await a message (OK, sid) or $(\text{Abort}, \text{sid})$ from the adversary; if the message is (OK, sid) then continue; otherwise, send the message $(\text{Abort}, \text{sid})$ to all honest parties and ignore all further messages to \mathcal{F}_{MPC} with this sid .

(continued...)

Fig. 3. Functionality $\mathcal{F}_{\text{Prep}}$

Functionality $\mathcal{F}_{\text{Prep}}$ (continued)

Additional commands

daBits On receiving $(\text{daBits}, \text{id}_1, \dots, \text{id}_\ell, \text{sid}_1, \text{sid}_2)$, from all parties where $\text{id}_i \notin \text{Val.Keys}$ for all $i \in [\ell]$, await a message **OK** or **Abort** from the adversary. If the message is **OK**, then sample $\{b_j\}_{j \in [\ell]} \xleftarrow{\$} \{0, 1\}$ and for each $j \in [\ell]$, set $\text{Val}_{\text{sid}_1}[\text{id}_j] \leftarrow b_j$ and $\text{Val}_{\text{sid}_2}[\text{id}_j] \leftarrow b_j$ and insert the set $\{\text{id}_i\}_{i \in [\ell]}$ into $\text{Val}_{\text{sid}_1}.\text{Keys}$ and $\text{Val}_{\text{sid}_2}.\text{Keys}$; otherwise send the messages $(\text{Abort}, \text{sid}_1)$ and $(\text{Abort}, \text{sid}_2)$ to all honest parties and the adversary and ignore all further messages to \mathcal{F}_{MPC} with session identifier sid_1 or sid_2 .

Fig. 4. Functionality $\mathcal{F}_{\text{Prep}}$ (continued)

4.1 Generating daBits using Bucketing

Any technique for generating daBits require some form of checking procedure to ensure consistency between the two fields. Checking consistency often means checking random linear combinations of secrets produce the same result in both cases. Unfortunately, in our case such comparisons are meaningless since the fields have different characteristic and it seems hard to find a mapping between the two fields which allows to compare random values in \mathbb{F}_p with values in \mathbb{F}_{2^k} . We can, however, check XORs of bits, which in \mathbb{F}_p involves multiplication. (See Equation 1 in Section 2.) It is therefore necessary to use a protocol that minimises (as far as possible) the number of multiplications. Consequently, techniques using oblivious transfer (OT) such as [WRK17] to generate authenticated bits require a lot of XORs for checking correctness, so are undesirable for generating daBits.

Our chosen solution uses \mathcal{F}_{MPC} as a black box. In order to generate the same bit in both fields, each party samples a bit and calls the \mathbb{F}_p and \mathbb{F}_{2^k} instances of \mathcal{F}_{MPC} with this same input and then the parties compute the n -party XOR. To ensure all parties provided the same inputs in both fields, cut-and-choose and bucketing procedures are required, though since the number of bits necessary to generate is a multiple of $\log p \cdot \text{sec}$ and we can batch-produce daBits, the parameters are modest (for $\text{sec} = 40$ five XORs in $\mathcal{F}_{\text{MPC}}^p$ per daBit are enough).

We use similar cut-and-choose and bucketing checks to those described by Frederiksen et al. [FKOS15, App. F.3], in which “triple-like” secrets can be efficiently checked. The idea behind these checks is the following. One first opens a random subset of secrets so that with some probability all unopened bits are correct. This ensures that the adversary cannot cheat on too many of the daBits. One then puts the secrets into buckets, and then in each bucket designates one secret as the one to output, uses all other secrets in the bucket to check the last, and discards all but the designated secret. For a single bucket, the check will only pass (by construction) if either all secrets are correct or all are incorrect. Thus the adversary is forced to corrupt whole multiples of the bucket size and hope they are grouped together in the same bucket. Fortunately, (we will show that) there is no leakage on the bits since the parameters required for the parts of the protocol described above already preclude it. The protocol is described in Figures 5 and 6; we prove that this protocol securely realises

the functionality $\mathcal{F}_{\text{Prep}}$ in Figures 3 and 4 in the \mathcal{F}_{MPC} -hybrid model. To do this, we require Proposition 4.1.

Protocol $\Pi_{\text{daBits+MPC}}$

This protocol is in the \mathcal{F}_{MPC} -hybrid model.

Initialise

1. Call an instance, $\mathcal{F}_{\text{MPC}}^p$, of \mathcal{F}_{MPC} with input (Initialise, $\mathbb{F}_p, 0$).
2. Call an instance, $\mathcal{F}_{\text{MPC}}^{2^k}$, of \mathcal{F}_{MPC} with input (Initialise, $\mathbb{F}_{2^k}, 1$).

Calls to \mathcal{F}_{MPC} Dealt with by $\mathcal{F}_{\text{MPC}}^p$ or $\mathcal{F}_{\text{MPC}}^{2^k}$, as appropriate.

daBits To generate ℓ bits, do the following:

1. Generate daBits

- (a) Let $m \leftarrow CB\ell$ where $C > 1$ and $B > 1$ are chosen so that $C^B \cdot \binom{B\ell}{B} > 2^{\text{sec}}$.
- (b) For each $i \in [n]$,
 - i. Party P_i samples a bit string $(b_1^i, \dots, b_m^i) \xleftarrow{\$} \{0, 1\}^m$.
 - ii. Call $\mathcal{F}_{\text{MPC}}^p$ where P_i has input (Input, $i, \text{id}_{b_j^i}, b_1^i, 0$) $_{j=1}^m$ and P_j ($j \neq i$) has input (Input, $i, \text{id}_{b_j^i}, \perp, 0$) $_{i=1}^m$.
 - iii. Call $\mathcal{F}_{\text{MPC}}^{2^k}$ where P_i has input (Input, $i, \text{id}_{b_j^i}, b_1^i, 1$) $_{j=1}^m$ and P_j ($j \neq i$) has input (Input, $i, \text{id}_{b_j^i}, \perp, 1$) $_{i=1}^m$.

2. Cut and Choose

- (a) Call $\mathcal{F}_{\text{Rand}}$ with input (RSubset, $[CB\ell], (C-1)B\ell$) to obtain a set S .
- (b) Call $\mathcal{F}_{\text{MPC}}^p$ with inputs (Open, 0, $\text{id}_{b_j^i}, 0$) $_{j \in S}$ for all $i \in [n]$.
- (c) Call $\mathcal{F}_{\text{MPC}}^{2^k}$ with inputs (Open, 0, $\text{id}_{b_j^i}, 1$) $_{j \in S}$ for all $i \in [n]$.
- (d) If any party sees daBits which are not in $\{0, 1\}$ or not the same in both fields, they send the message **Abort** to all parties and halt.

3. Combine

For all $j \in S$, do the following:

- (a) Set $\llbracket b_j \rrbracket_p \leftarrow \llbracket b_j^1 \rrbracket_p$ and then for i from 2 to n compute
 - i. $\llbracket b_j \rrbracket_p \leftarrow \llbracket b_j \rrbracket_p + \llbracket b_j^i \rrbracket_p - 2 \cdot \llbracket b_j \rrbracket_p \cdot \llbracket b_j^i \rrbracket_p$
- (b) Compute $\llbracket b_j \rrbracket_{2^k} \leftarrow \bigoplus_{i=1}^n \llbracket b_j^i \rrbracket_{2^k}$.

(continued...)

Fig. 5. Protocol $\Pi_{\text{daBits+MPC}}$

Protocol $\Pi_{\text{daBits+MPC}}$

4. Consistency Check

- (a) Call $\mathcal{F}_{\text{Rand}}$ with input $(\text{RBucket}, [B\ell], B)$ and use the returned sets $(S_i)_{i=1}^\ell$ to put the $B\ell$ daBits into ℓ buckets of size B .
- (b) For each bucket S_i ,
 - i. Relabel the bits in this bucket as b^1, \dots, b^B .
 - ii. For $j = 2$ to B , compute $\llbracket c^j \rrbracket_p \leftarrow \llbracket b^1 \rrbracket_p + \llbracket b^j \rrbracket_p - 2 \cdot \llbracket b^1 \rrbracket_p \cdot \llbracket b^j \rrbracket_p$ and $\llbracket c^j \rrbracket_{2^k} \leftarrow \llbracket b^1 \rrbracket_{2^k} \oplus \llbracket b^j \rrbracket_{2^k}$.
 - iii. Call $\mathcal{F}_{\text{MPC}}^p$ with inputs $(\text{Open}, 0, \text{id}_{c^j}, 0)_{j=2}^B$.
 - iv. Call $\mathcal{F}_{\text{MPC}}^{2^k}$ with inputs $(\text{Open}, 0, \text{id}_{c^j}, 1)_{j=2}^B$.
 - v. If any party sees daBits which are not in $\{0, 1\}$ or not the same in both fields, they send the message **Abort** to all parties and halt.
 - vi. Set $\llbracket b_i \rrbracket_{p, 2^k} \leftarrow \llbracket b^1 \rrbracket_{p, 2^k}$.
- (c) Call $\mathcal{F}_{\text{MPC}}^p$ with input $(\text{Check}, 0)$.
- (d) Call $\mathcal{F}_{\text{MPC}}^{2^k}$ with input $(\text{Check}, 1)$.
- (e) If the checks pass without aborting, output $\{\llbracket b_i \rrbracket_{p, 2^k}\}_{i=1}^\ell$ and discard all other bits.

Fig. 6. Protocol $\Pi_{\text{daBits+MPC}}$ (continued)

Proposition 4.1. *For a given $\ell > 0$, choose $B > 1$ and $C > 1$ so that $C^{-B} \cdot \binom{B\ell}{B}^{-1} < 2^{-\text{sec}}$. Then the probability that one or more of the ℓ daBits output after **Consistency Check** by $\Pi_{\text{daBits+MPC}}$ in Figure 6 is different in each field is at most $2^{-\text{sec}}$.*

Proof. Using $\mathcal{F}_{\text{MPC}}^p$ and $\mathcal{F}_{\text{MPC}}^{2^k}$ as black boxes ensures the adversary can only possibly cheat in the input stage. We will argue that:

1. If both sets of inputs from corrupt parties to $\mathcal{F}_{\text{MPC}}^p$ and $\mathcal{F}_{\text{MPC}}^{2^k}$ are bits (rather than other field elements), then the bits are consistent in the two different fields with overwhelming probability.
2. The inputs in \mathbb{F}_{2^k} are bits with overwhelming probability.
3. The inputs in \mathbb{F}_p are bits with overwhelming probability.

We will conclude that the daBits are bits in the two fields, and are consistent.

1. Let c be the number of inconsistent daBits generated by a given corrupt party. If $c > B\ell$ then every set of size $(C-1)B\ell$ contains an incorrect daBit so the honest parties will always detect this in **Cut and Choose** and abort. Since $(C-1)B\ell$ out of $CB\ell$ daBits are opened, on average the probability that a daBit is not opened is $1 - (C-1)/C = C^{-1}$, and so if $c < B\ell$ then we have:

$$\Pr[\text{None of the } c \text{ corrupted daBits is opened}] = C^{-c}. \quad (2)$$

At this point, if the protocol has not yet aborted, then there are $B\ell$ daBits remaining of which exactly c are corrupt.

Suppose a daBit $\llbracket b \rrbracket_{p, 2^k}$ takes the value \tilde{b} in \mathbb{F}_p and \hat{b} in \mathbb{F}_{2^k} . If the bucketing check passes then for every other daBit $\llbracket b' \rrbracket_{p, 2^k}$ in the bucket it holds that $\tilde{b} \oplus \tilde{b}' = \hat{b} \oplus \hat{b}'$, so $\tilde{b}' = (\hat{b} \oplus \hat{b}') \oplus \tilde{b}$,

and so $\tilde{b} = \hat{b} \oplus 1$ if and only if $\tilde{b}' = \hat{b}' \oplus 1$. (Recall that we are assuming the inputs are certainly bits at this stage.) In other words, within a single bucket, the check passes if and only if either all daBits are inconsistent, or if none of them are. Thus the probability **Consistency Check** passes without aborting is the probability that all corrupted daBits are placed into the same buckets. Moreover, this implies that if the number of corrupted daBits, c , is not a multiple of the bucket size, this stage never passes, so we write $c = Bt$ for some $t > 0$. Then we have:

$$\begin{aligned} \Pr[\text{All corrupted daBits are placed in the same buckets}] &= \\ &= \frac{\binom{Bt}{B} \cdot \binom{B(t-1)}{B} \cdots \binom{B}{B} \cdot \binom{B\ell-Bt}{B} \cdot \binom{B\ell-Bt-B}{B} \cdots \binom{B}{B}}{\binom{B\ell}{B} \cdot \binom{B\ell-B}{B} \cdots \binom{B}{B}} \\ &= \frac{(Bt)!}{B!^t} \cdot \frac{(B\ell-Bt)!}{B!^{\ell-t}} \cdot \frac{B!^\ell}{(B\ell)!} = \left(\frac{B\ell}{Bt}\right)^{-1}. \end{aligned} \quad (3)$$

Since the randomness for **Cut and Choose** and **Check Correctness** is independent, the event that both checks pass after the adversary corrupts c daBits is the product of the probabilities. To upper-bound the adversary's chance of winning, we compute the probability by maximising over t : thus we need C and B so that

$$\max_t \left\{ C^{-Bt} \cdot \left(\frac{B\ell}{Bt}\right)^{-1} \right\} < 2^{-\text{sec}} \quad (4)$$

The maximum occurs when t is small, and $t \geq 1$ otherwise no cheating occurred; thus since the proposition stipulates that $C^{-B} \cdot \left(\frac{B\ell}{B}\right)^{-1} < 2^{-\text{sec}}$, the daBits are consistent in both fields, if they are indeed bits in both fields.

2. Next, we will argue that the check in **Cut and Choose** ensures that the inputs given to $\mathcal{F}_{\text{MPC}}^{2^k}$ are indeed bits. It follows from Equation 2 that the step **Cut and Choose** aborts with probability C^{-c} if any element of either field is not a bit, as well as if the element in the two fields does not match. Moreover, in **Consistency Check**, in order for the check to pass in \mathbb{F}_{2^k} for a given bucket, the secrets' higher-order bits must be the same for all shares so that the XOR is always zero when the pairwise XORs are opened. Thus the probability that this happens is the same as the probability above in Equation 4 since again this can only happen when the adversary is not detected in **Cut and Choose**, that he cheats in some multiple of B daBits, and that these cheating bits are placed in the same buckets in **Consistency Check**.

3. We now show that all of the \mathbb{F}_p components are bits. To do this, we will show that if the \mathbb{F}_p component of a daBit is not a bit, then the bucket check passes only if all other daBits in the bucket are also not bits in \mathbb{F}_p .

If the protocol has not aborted, then in every bucket B , for every $2 \leq j \leq B$, it holds that

$$b^1 + b^j - 2 \cdot b^1 \cdot b^j = c^j \quad (5)$$

where $c^j \in \{0, 1\}$ are determined by the XOR in \mathbb{F}_{2^k} . Note that since $c^j = \bigoplus_{i=1}^n b_i^1 \oplus \bigoplus_{i=1}^n b_i^j$ and at least one b_i^j is generated by an honest party, this value is uniform and unknown to the adversary when he chooses his inputs at the beginning.

Suppose $b^1 \in \mathbb{F}_p \setminus \{0, 1\}$. If $b^1 = 2^{-1} \in \mathbb{F}_p$ then by Equation 5 we have $b^1 = c^j$; but c^j is a bit, so the “XOR” is not the same in both fields and the protocol will abort. Thus we may assume $b^1 \neq 2^{-1}$ and so we can rewrite the equation above as

$$b^j = \frac{b^1 - c^j}{2 \cdot b^1 - 1}. \quad (6)$$

Now if b^j is a bit then it satisfies $b^j(b^j - 1) = 0$, and so

$$0 = \left(\frac{b^1 - c^j}{2 \cdot b^1 - 1} \right) \cdot \left(\frac{b^1 - c^j}{2 \cdot b^1 - 1} - 1 \right) = -\frac{(b^1 - c^j)(b^1 - (1 - c^j))}{(2 \cdot b^1 - 1)^2}$$

so $b^1 = c^j$ or $b^1 = 1 - c^j$; thus $b^1 \in \{0, 1\}$, which is a contradiction. Thus we have shown that if b^1 is not a bit then b^j is not a bit for every other b^j in this bucket. Moreover, for each $j = 2, \dots, B$, there are two distinct values $b^j \in \mathbb{F}_p \setminus \{0, 1\}$ solving Equation 6 corresponding to the two possible values of $c^j \in \{0, 1\}$, which means that if the bucket check passes then the adversary must *also* have guessed the bits $\{c^j\}_{j=1}^B$, which he can do with probability 2^{-B} since they are constructed using at least one honest party’s input. Thus the chance of cheating without detection in this way is at most $2^{-Bt} \cdot C^{-Bt} \cdot \binom{Bt}{B}^{-1}$.

Thus we have shown that the probability that $b^1 \in \mathbb{F}_p \setminus \{0, 1\}$ is given as output for the \mathbb{F}_p component is at most the probability that the adversary corrupts a multiple of B daBits, that these daBits are placed in the same buckets, and that the adversary correctly guesses c bits from honest parties (in the construction of the bits $\{b^j\}_{j \in [B]}$) so that the appropriate equations hold in the corrupted buckets. Indeed, needing to guess the bits ahead of time only *reduces* the adversary’s chance of winning from the same probability in the \mathbb{F}_{2^k} case.

We conclude that the daBits are bits in both fields and are the same in both fields except with probability at most $2^{-\text{sec}}$. \square

Theorem 4.1. *The protocol $\Pi_{\text{daBits}+\text{MPC}}$ securely realises $\mathcal{F}_{\text{Prep}}$ in the $(\mathcal{F}_{\text{MPC}}, \mathcal{F}_{\text{Rand}})$ -hybrid model against an active adversary corrupting up to $n - 1$ out of n parties.*

Proof. To prove security in the UC framework we must show that to any environment \mathcal{Z} , for any adversary \mathcal{A} there exists a simulator \mathcal{S} such that the execution of an idealised version of the protocol run by a trusted third party \mathcal{F} with the simulator is indistinguishable from a real execution of the protocol Π between the honest parties and the adversary. The environment specifies the code run by the adversary as well as the inputs of all parties, honest and dishonest. Additionally, the environment sees all outputs of all parties; it does not see the intermediate interactions in subroutines of the honest parties’ executions, otherwise distinguishing would be trivial as honest parties either perform Π or interact with \mathcal{F} . In the $(\mathcal{F}_{\text{MPC}}, \mathcal{F}_{\text{Rand}})$ -hybrid model, the adversary is allowed to make oracle queries to these functionalities and \mathcal{S} must generate the responses.

Note the functionality does *not* have access to the random tapes honest parties as this would make distinguishing between worlds trivial: it would be impossible for the simulator to emulate honest parties to the real-world adversary indistinguishably since for any random tape sampled by the simulator, the environment would always be able to execute the protocol internally, using its knowledge of the random tapes of honest parties to execute the entire protocol deterministically, and compare it to the output of the simulator.

Following standard practice, and as described in [Can00, §4.2.2], we define a simulator which interacts with the adversary \mathcal{A} as a black box. This allows us to make the claim that the simulator works regardless of the code run by the adversary and hence prove the claim.

Suppose the adversary corrupts $t < n$ parties in total, indexed by a set A . We define a sequence of hybrid worlds (**Hybrid** h) $_{h=0}^{n-t}$ and show that each is indistinguishable from the previous. **Hybrid** h is defined as:

Hybrid h The simulator has the actual input of $n - t - h$ honest parties and must simulate the remaining h honest parties towards the adversary.

The simulator is described in Figure 7.

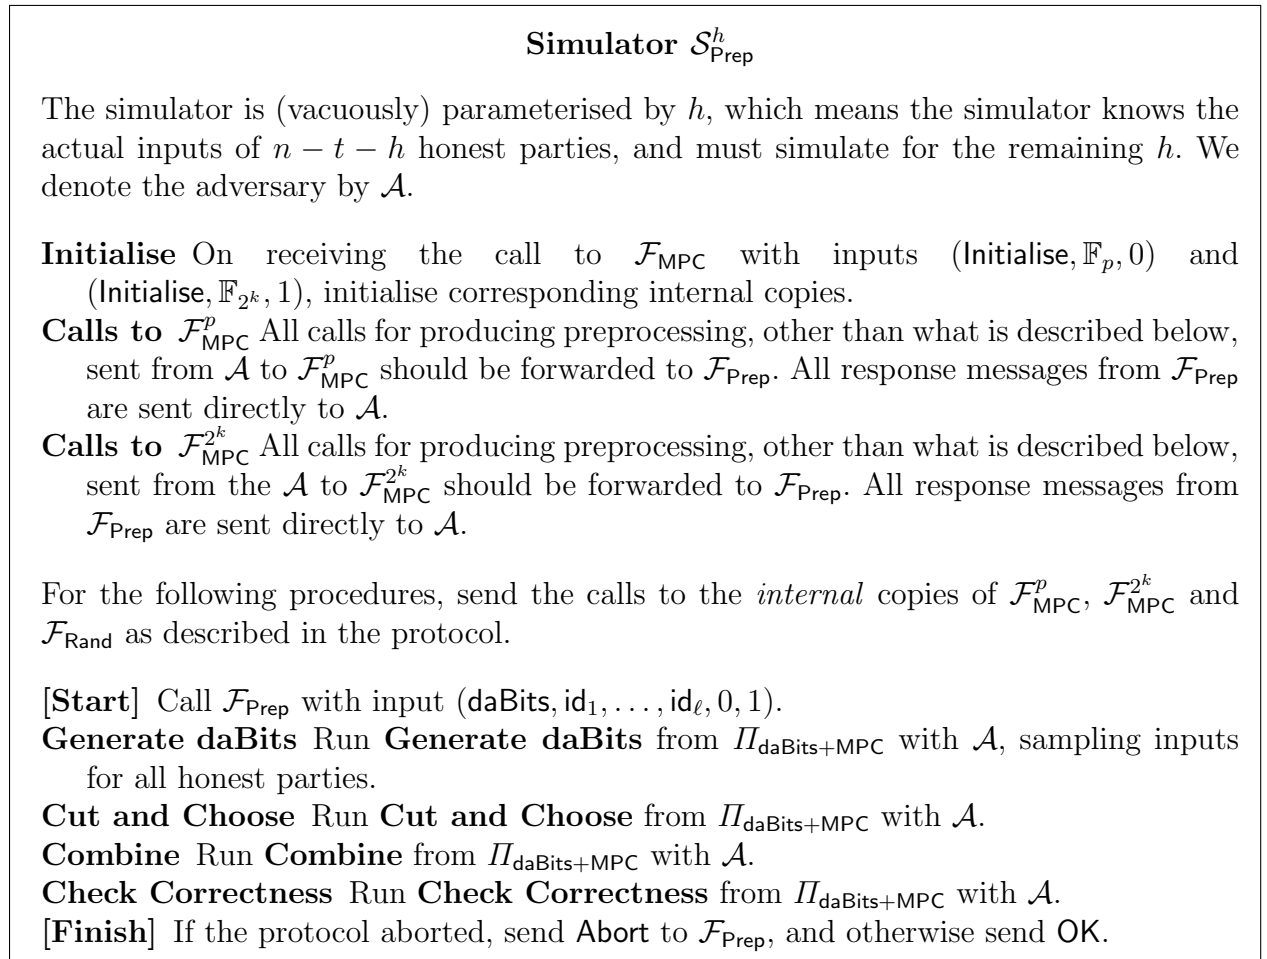


Fig. 7. Simulator $\mathcal{S}_{\text{Prep}}^h$

Claim. The $\mathcal{F}_{\text{MPC}}, \mathcal{F}_{\text{Rand}}$ -hybrid world is indistinguishable from **Hybrid 0**.

Proof. Correctness of the simulation holds as follows. The simulator emulates $\mathcal{F}_{\text{MPC}}^p$, $\mathcal{F}_{\text{MPC}}^{2^k}$ and $\mathcal{F}_{\text{Rand}}$, so all calls made to these oracles are dealt with as in an execution of the protocol. Indeed, for all calls to \mathcal{F}_{MPC} in either field which are outside of the daBits generation procedure, the commands are forwarded to $\mathcal{F}_{\text{Prep}}$ and relayed back to \mathcal{A} , and since $\mathcal{F}_{\text{Prep}}$ has the same interface as \mathcal{F}_{MPC} by definition, there is no difference between the worlds. As for the daBit generation, when the adversary makes calls to provide (random) inputs and then perform **Cut and Choose**, the simulator does not forward the messages through to $\mathcal{F}_{\text{Prep}}$ since all bits used in the protocol except the final output bits are discarded. Instead the command $(\text{daBits}, \text{id}_1, \dots, \text{id}_\ell, 0, 1)$ is sent to $\mathcal{F}_{\text{Prep}}$ and the simulator executes the daBit routines honestly with the adversary, making random choices for honest parties by sampling in the same way as in the protocol.

Now we argue indistinguishability between executions: we must show that for any algorithm \mathcal{A} specified by the environment \mathcal{Z} , it holds that

$$\text{EXEC}(\mathcal{Z}, \mathcal{A}^{\mathcal{F}_{\text{MPC}}, \mathcal{F}_{\text{Rand}}}, \Pi_{\text{daBits}+\text{MPC}}) \sim \text{EXEC}(\mathcal{Z}, \mathcal{S}_{\text{Prep}}^0, \mathcal{F}_{\text{Prep}})$$

where \sim denotes statistical indistinguishability of distributions, and the randomness of these distributions is taken over the random tapes of honest parties and the adversary and simulator.

First, note that the oracles \mathcal{F}_{MPC} and $\mathcal{F}_{\text{Rand}}$ are executed honestly by $\mathcal{S}_{\text{Prep}}^0$ so the contribution to the distributions is the same in both executions.

Second, since the inputs of honest parties are sampled during the protocol, they are not specified or known by the environment. However, if the adversary performs a *selective-failure attack*, then the environment may learn information. A selective failure attack is where the environment can learn some information if the protocol does not detect cheating behaviour. For example, if the environment guesses an entire bucket of bits and chooses inputs for the adversary's input so that the bucket check would pass based on these guesses, then if the protocol does not abort then the environment learns that its guesses were correct. Then if the final output bit is *not* the XOR of all parties' inputs then the execution must have happened in **Hybrid 0** since in this world the output depends on the random tape of $\mathcal{F}_{\text{Prep}}$ and is independent of the adversary's and honest parties' random tapes, contrasting the output in the $\mathcal{F}_{\text{MPC}}, \mathcal{F}_{\text{Rand}}$ -hybrid world in which the final output is an XOR of bits on these tapes (which were guessed by the environment). Since this happens with probability $\frac{1}{2}$, in expected 2 executions, the environment can distinguish. However, by Proposition 4.1, the environment can only mount a selective failure attack with success with probability at most $2^{-\text{sec}}$ by the choice of parameters.

Thus the only way to distinguish between worlds is if the transcript leaks information on the honest parties' inputs. In **Check Correctness**, XORs are computed in both fields and the result is opened; however, this reveals no information on the final daBit outputs as the linear dependence between the secret and the public values is broken by discarding all secrets in each bucket except the designated (i.e. first) bit. We conclude that the overall distributions of the two executions are statistically indistinguishable in **sec**. ■

Claim. **Hybrid** h is indistinguishable from **Hybrid** $h + 1$ for $h = 0, \dots, n - t - 1$.

Proof. There is no difference between these worlds since honest parties' (random) inputs are sampled the same way in both cases. ■

Since \mathcal{F}_{MPC} is secure up to $t = n - 1$, the result follows. □

4.2 Garbling and Switching

In this section we give a high-level description of how our approach can be used to provide input to a garbled circuit from secret-shared data, and convert garbled-circuit outputs into sharings of secrets in \mathbb{F}_p .

In Figure 8 we show pictographically what happens at the barrier between secret-sharing and garbling, though note we have shown a circuit output that will be reconstructed to an element of \mathbb{F}_p : the circuit output can be any string of bits. As was discussed in the preliminaries, many recent garbling protocols use specialised secret-sharing MPC protocols. On this basis, it is straightforward to see that from $\mathcal{F}_{\text{Prep}}$ the parties can both perform secret-sharing-based MPC over a large prime finite field and can garble circuits. A concrete example of this is given in Figures 16, 17, 18 and 19, which modifies the SPDZ-BMR protocol [LPSY15] to allow for our conversion.




Fig. 8. Overview

From SS to GC In brief, the parties input a secret-shared $\llbracket x \rrbracket_p$ by computing $\llbracket x - r \rrbracket_p$ and opening it to reveal $x - r$ where $r = \sum_{j=0}^{\lceil \log p \rceil - 1} 2^j \cdot \llbracket r_j \rrbracket_p$ is constructed from $\text{daBits} \{ \llbracket r_j \rrbracket_{p, 2^k} \}_{j=0}^{\lceil \log p \rceil - 1}$, and \mathcal{F}_{MPC} is called with input $(\text{Check}, 0)$ either at this point or later on, and then these public values are taken to be input bits to the garbled circuit. To correct the offset r , the circuit $(x - r) + r \bmod p$ is computed inside the garbled circuit. This is possible since the bits of r can be hard-wired into the circuit using the \mathbb{F}_{2^k} sharings of its bit-decomposition.

Note that typically for a party to provide input bit b on wire w in a garbled circuit, the parties reveal the secret-shared *wire mask* $\llbracket \lambda_w \rrbracket_{2^k}$ to this party, which broadcasts $\Lambda_w \leftarrow b \oplus \lambda_w$, called the associated *signal bit*; then the parties communicate further to reveal keys required for ciphertext decryptions, which is how the circuit is evaluated. This mask thus hides the actual input (and is removed inside the garbled circuit). Since the inputs here are the bits

of the public value $x - r$, there is no need mask inputs here, and thus it suffices to set all the corresponding wire mask bits to be 0.

From GC to SS In standard BMR-style garbling protocols, the outputs of the circuit are a set of public signal bits. These are equal to the actual Boolean outputs XORed with circuit output wire masks, which are initially secret-shared, concealing the actual outputs. Typically in multi-party circuit garbling, the wire masks for output wires are revealed immediately after the garbling stage so that all parties can learn the final outputs without communication after locally evaluating the garbled circuit. When garbling circuits using SS-based techniques, and aiming for computation in which parties can continue to operate on private outputs of a GC, a simple way of obtaining shared output is for the parties not to reveal the secret-shared wire masks for output wires after garbling and instead, after evaluating, to compute the XOR of the secret-shared mask with the public signal bit, in MPC.

In other words, for output wire w they obtain a sharing of the secret output bit b by computing

$$\llbracket b \rrbracket_{2^k} \leftarrow \Lambda_w \oplus \llbracket \lambda_w \rrbracket_{2^k}.$$

In our case, we want the shared output of the circuit to be in \mathbb{F}_p , and to do this it suffices for the masks on circuit output wires to be daBits (instead of random bits shared only in \mathbb{F}_{2^k} as would be done normally) and for the parties to compute (locally)

$$\llbracket b \rrbracket_p \leftarrow \Lambda_w + \llbracket \lambda_w \rrbracket_p - 2 \cdot \Lambda_w \cdot \llbracket \lambda_w \rrbracket_p.$$

To avoid interfering with the description of the garbling subprotocol, we can define an additional layer to the circuit after the output layer which converts output wires with masks only in \mathbb{F}_{2^k} to output wires with masks as daBits, without changing the real values on the wire. To do this, for every output wire w , let $\llbracket \lambda_w \rrbracket_{2^k}$ be the associated secret-shared wire mask. Then,

- In the garbling stage take a new daBit $\llbracket \lambda_{w'} \rrbracket_{p,2^k}$,
 1. Set $\llbracket \Lambda_{w_0} \rrbracket_{2^k} \leftarrow \llbracket \lambda_w \rrbracket_{2^k} \oplus \llbracket \lambda_{w'} \rrbracket_{2^k}$.
 2. Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, 0, \text{id}_{\Lambda_{w_0}}, 1)$ to obtain Λ_{w_0} .
- In the evaluation stage, upon obtaining Λ_w ,
 1. Compute $\Lambda_{w'} \leftarrow \Lambda_w \oplus \Lambda_{w_0}$.
 2. Compute the final (\mathbb{F}_p -secret-shared) output as $\llbracket b \rrbracket_p \leftarrow \Lambda_{w'} + \llbracket \lambda_{w'} \rrbracket_p - 2 \cdot \Lambda_{w'} \cdot \llbracket \lambda_{w'} \rrbracket_p$.

Observe that $\Lambda_{w_0} \equiv \lambda_{w_0}$ so this procedure is just adding a layer of XOR gates where the masking bits are daBits and the other input wire is always 0 (so the gate evaluation doesn't change the real wire value). Note that since the signal bits for XOR gates are determined from input signal bits and not the output key, there is no need to generate an output key for wire w_0 . For correctness, observe that




Fig. 9. Circuit output wires

$$\begin{aligned}
\Lambda_{w'} \oplus \lambda_{w'} &= (\Lambda_w \oplus \Lambda_{w_0}) \oplus (\lambda_w \oplus \lambda_{w_0}) \\
&= ((b \oplus \lambda_w) \oplus (0 \oplus \lambda_{w_0})) \oplus (\lambda_w \oplus \lambda_{w_0}) \\
&= b.
\end{aligned}$$

Correctness of the actual garbling was outlined in Section 2. The proof of Theorem C.1 is deferred to the appendices as it is straightforward, following from the security of SPDZ-BMR [LPSY15] and the fact that the additional input and output procedures perfectly hide the actual circuit inputs and outputs.

5 Implementation

We have implemented daBit generation and the conversion between arithmetic shares and garbled circuits. Our code is developed on top of the MP-SPDZ framework [Ana19] and experiments were run on computers with an identical configuration of i7-7700K CPU and 32GB RAM connected via a 1Gb/s LAN connection with an average round-trip ping time of 0.3ms. The $\mathcal{F}_{\text{MPC}}^p$ functionality is implemented using LowGear, one of the two variants of Overdrive [KPR18], a SPDZ-like protocol; the $\mathcal{F}_{\text{MPC}}^{2^k}$ functionality is implemented using MASCOT [KOS16] to realise the protocol of [KY18], a variant of SPDZ-BMR multiparty circuit garbling [LPSY15]. In our experiments, \mathbb{F}_{2^k} is always taken with $k = \kappa = 128$ since this is the security of PRF keys used in SPDZ-BMR. The daBits are always generated with $\kappa = 128$ and the same statistical security sec as the protocol for \mathcal{F}_{MPC} .

Primes. We require that p be close to a power of 2 so that $x - r$ is indistinguishable from a uniform element of the field, as discussed in Section 2. Since we use LowGear in our implementation, for a technical reason we also require that p be congruent to 1 mod N where $N = 32768$. (This is the amount of packing in the ciphertexts.) Consequently, using LowGear means we always lose $15 = \log 32768$ bits of statistical security if $p > 65537$ since then the k -bit prime must be of the form $2^{k-1} + t \cdot 2^{15} + 1$ for some t where $1 \leq t \leq 2^{k-16} - 1$. We stress that the loss of 15 bits (out of $\log p$) statistical security can be mitigated by executing the conversion only after a check of $r < p$ inside the GC. In this way we guarantee that r is a truly random element modulo p . If we settle for bounded inputs $\log |x| < \log p - \text{sec}$ then the check can be skipped as $x - r$ is close to uniform with distance $2^{-\text{sec}}$ where $r = \sum_{i=0}^{\log |x| + \text{sec}} r_i$.

Cut and choose optimisation. One key observation that enables reduction of the preprocessing overhead in \mathbb{F}_{2^k} is that parties only need to input bits (instead of full \mathbb{F}_{2^k} field elements) into \mathcal{F}_{MPC} during $\Pi_{\text{daBits}+\text{MPC}}$. For a party to input a secret x in MASCOT, the parties create a random authenticated mask r and open it to the party, and the party then broadcasts $x + r$. Since the inputs are just bits, it suffices for the random masks also to be bits. Generating authenticated bits using MASCOT is extremely cheap and comes with a small communication overhead (see Table 2).

More efficient packing for MAC Check. Instead of a set of k secret bits being opened as full \mathbb{F}_{2^k} field elements $(0, \dots, 0, b_1), \dots, (0, \dots, 0, b_t) \in \mathbb{F}_2^k \cong \mathbb{F}_{2^k}$, we can save on all the redundant

0's being sent by sending a single field element $(b_k, \dots, b_1) \in \mathbb{F}_{2^k}$. This optimisation reduces by a factor 2 the amount of data sent for the online phase of daBit generation.

Complexity analysis. In LowGear (Overdrive) and MASCOT the authors avoided reporting any benchmarks for random bit masks in \mathbb{F}_{2^k} or random input masks in \mathbb{F}_p since they focused on the entire triple generation protocol. Fortunately their code is open source and easy to modify so we micro-benchmarked their protocols in order to get concrete costs for the procedure **Input** for $\mathcal{F}_{\text{MPC}}^p$ and $\mathcal{F}_{\text{MPC}}^{2^k}$. For example, in the two-party case, to provide a bit as input bit costs overall 0.384kb with MASCOT in \mathbb{F}_{2^k} , whereas for LowGear, providing bits as input costs the same as any arbitrary field element in \mathbb{F}_p , requiring 2.048kb. Hence, with the current state of protocols, inputs are cheap in a binary field whereas triples are cheap in a prime field.




Fig. 10. Total communication costs for all parties per preprocessed element.

Bucketing parameters. Recall that our goal is to minimise the total amount of communication and time spent by parties generating each daBit. By examining the input and triple costs for LowGear and MASCOT (see Table 6 in Appendix E) the optimal communication for statistical security $\text{sec} = 64$ and $p \approx 2^{128}$ was found to occur when generating $l = 8192$ daBits per loop, a cut-and-choose parameter $C = 5$ and a bucket size $B = 4$. Then we ran the daBit generation along with LowGear and MASCOT for multiple parties on the same computers configuration to get the total communication cost in order to see how the costs scale with the number of parties. Results are given in Figure 10. While MASCOT triples are not used during the daBit production, we believe that comparing the cost of a daBit to the best triple generation in \mathbb{F}_{2^k} helps to give a rough idea of how expensive a single daBit is.

To see how efficiency scales when the statistical security parameter sec is increased, we record the fewest numbers of calls to \mathcal{F}_{MPC} , optimising for total (actual) communication cost in Table 1. Since the numbers are dependent on integers (number of parties, size of buckets, and cut and choose parameter), several of the numbers in the table give far better security than the minimum stated. Note that since we optimise for the total communication cost and not for the smallest **Cut and Choose** and **Bucketing** parameters that achieve each level of security, in the cost for $\text{sec} = 64$ the number of calls to $\mathcal{F}_{\text{MPC}}.$ **Input** is larger than

# daBits	sec > 40			sec > 64			sec > 80		
	128	1024	8192	128	1024	8192	128	1024	8192
Calls to $\mathcal{F}_{\text{MPC}}^{\{p, 2^k\}}$. Input	40	16	12	42	40	40	36	28	24
Calls to $\mathcal{F}_{\text{MPC}}^p$. Multiply	7	7	5	13	9	7	17	13	11
Achieved sec	40	47	44	67	64	64	82	84	90

Table 1. Two parties pre-processing costs per daBit while varying the number of daBits per batch and statistical security. Parameters minimize for total communication given by LowGear and MASCOT.

for $\text{sec} = 80$. The bucket size, correlated with the number of calls to $\mathcal{F}_{\text{MPC}}.$ **Multiply**, is therefore is smaller than for $\text{sec} = 80$.

sec log p k	Comm. (kb)			Total (kb)	Time (ms)			Total(ms)
	$\mathcal{F}_{\text{MPC}}^p$	$\mathcal{F}_{\text{MPC}}^{2^k}$	daBitgen		$\mathcal{F}_{\text{MPC}}^p$	$\mathcal{F}_{\text{MPC}}^{2^k}$	daBitgen	
40 128 128	76.60	2.30	6.94	85.84	0.159	0.004	0.004	0.167
64 128 128	146.47	7.68	9.39	163.54	0.303	0.015	0.010	0.328
80 128 128	192.95	4.60	7.32	204.88	0.485	0.009	0.008	0.502

Table 2. 1Gb/s LAN experiments for two-party daBit generation per party. For all cases, the daBit batch has length 8192.

5.1 Switch between SPDZ (\mathbb{F}_p) to SPDZ-BMR (\mathbb{F}_{2^k})

To reduce the amount of garbling when converting an additive share to a GC one, we assume the \mathbb{F}_p input to the garbled circuit is bounded by $p/2^{\text{sec}}$. In this way then a uniform r in \mathbb{F}_p is 2^{sec} times larger than a so $x - r$ is statistically-indistinguishable from a uniform element of \mathbb{F}_p ; consequently, one need only garble $x + r$ and not $x + r \bmod p$, which makes the circuit marginally smaller – 379 AND gates for a 128 bit prime rather than ≈ 1000 AND gates for an addition mod p circuit.

In Table 3 we split the conversion into two phases: the total cost of generating 127 daBits for doing a full conversion (including the preprocessing triples from LowGear) and the online phase of SPDZ-BMR.

Comparison to semi-honest conversion. Keeping the same security parameters and computers configuration as [RWT⁺18, DSZ15], when benchmarked with $\text{sec} = 40$, the online phase to convert 1000 field elements of size 32 bits takes 193ms. Our solution benefits from merging multiple conversions at once due to the SIMD nature of operations and that we can perform a single MAC-Check to compute the signal bits for the GC. Note that our conversion from an arithmetic SPDZ share to a SPDZ-BMR GC share takes about 14 times more than the semi-honest arithmetic to an Yao GC conversion in ABY or Chameleon [RWT⁺18, DSZ15].

Conversion	daBit (total)		SPDZ-BMR	
	Comm. (kbits)	Time (ms)	ANDs	Online (ms)
SPDZ \mapsto GC	20769	39.751	379	0.106
GC \mapsto SPDZ	10303	19.719	0	0.005

Table 3. Two parties 1Gb/s LAN experiments converting a 63 bit field element with 64 statistical security. BMR online phase times are amortized over 1000 executions in parallel (single-threaded).

5.2 Multiple class Support Vector Machine

A support vector machine (SVM) is a machine learning algorithm that uses training data to compute a matrix A and a vector \mathbf{b} such that for a so-called *feature vector* \mathbf{x} of a new input, the index of the largest component of the vector $A \cdot \mathbf{x} + \mathbf{b}$ is defined to be its class. We decided to benchmark this circuit using actively-secure circuit marbling as it is clear that there is an operation best suited to arithmetic circuits (namely, $\llbracket A \rrbracket \cdot \llbracket \mathbf{x} \rrbracket + \llbracket \mathbf{b} \rrbracket$) and another better for a Boolean circuit (namely, **argmax**, which computes the index of the vector’s largest component).

We have benchmarked the online phase of a multi-class Linear SVM with 102 classes and 128 features over a simulated WAN network (using the Linux `tc` command) with a round-trip ping time of 100ms and 50Mb/s bandwidth with two parties. The SVM structure is the same used by Makri et al. [MRSV19] to classify the Caltech-101 dataset which contains 102 different categories of images such as aeroplanes, dolphins, helicopters and others [FFFP04]. In this dataset, $\mathbf{x} \in \mathbb{F}_p^{128}$, $A \in \mathbb{F}_p^{102 \times 128}$ and $\mathbf{b} \in \mathbb{F}_p^{102}$, and it requires 102 conversions from \mathbb{F}_p to \mathbb{F}_2^k – one for each SVM label. The particular SVM used by Makri et al. has bounded inputs x where $\log |x| \leq 25$, a field size $\log p = 128$ and statistical security $\text{sec} = 64$.

We have implemented a special instruction in MP-SPDZ which loads a secret integer modulo p (a SPDZ share) into the SPDZ-BMR machine. To merge all modulo p instructions of SPDZ shares into SPDZ-BMR to form an universal Virtual Machine requires some extra engineering effort: this is why we chose to micro-benchmark in Table 4 the different stages of the online phase: doing $\llbracket \mathbf{y} \rrbracket_p \leftarrow \llbracket A \rrbracket_p \cdot \llbracket \mathbf{x} \rrbracket_p + \llbracket \mathbf{b} \rrbracket_p$ with SPDZ, then the instruction converting $\llbracket \mathbf{y} \rrbracket_p = (\llbracket y_1 \rrbracket_p, \dots, \llbracket y_{102} \rrbracket_p)$ to $(\{\llbracket (y_1)_j \rrbracket_{2^k} \}_{j=0}^{\log p-1}, \dots, \{\llbracket (y_{102})_j \rrbracket_{2^k} \}_{j=0}^{\log p-1})$, ending with the evaluation stage of SPDZ-BMR on

$$\text{argmax}((\llbracket (y_1)_j \rrbracket_{2^k} \}_{j=0}^{\log p-1}, \dots, (\llbracket (y_{102})_j \rrbracket_{2^k} \}_{j=0}^{\log p-1})).$$

We name this construction Marbled-SPDZ.

Online cost. The online phase (Table 4) using Marbled-SPDZ is more than 10 times faster than SPDZ-BMR and about 10 times faster than SPDZ.

Preprocessing cost. The preprocessing effort for the garbling (in AND gates) is reduced by a factor of almost 400 times using our construction. We chose to express the preprocessing costs of Table 4 in terms of AND gates, random triples and bits mainly for the reason that SPDZ-BMR requires much more work for an AND gate than WRK. Based on the concrete

preprocessing costs we have in Table 4 we give estimations on the communication where the preprocessing of the garbling is done via WRK: performing an SVM evaluation using i) WRK alone would require 6.6GB sent per party (3.8kb per AND gate), ii) SPDZ alone (with LowGear) would require 54MB per party (15kb per triple/random bit), iii) Marbled-SPDZ would take 160MB per party.

Nevertheless, the main cost in Marbled SPDZ is the daBit generation (119 MB) which is more than 70% of the preprocessing effort. If one chooses $\text{sec} = 40$ then we need five triples per daBit and 65 daBits per conversion which amounts to only 112.3MB for the entire SVM evaluation (twice the cost of plain SPDZ). A detailed cost can be found in Table 5.

Protocol	Online cost			Preprocessing cost		
	Comm. rounds	Time (ms)	Total (ms)	\mathbb{F}_p triples	\mathbb{F}_p bits	AND gates
SPDZ	54	2661	2661	19015	9797	-
SPDZ-BMR	0	2786	2786	-	-	14088217
Marbled-SPDZ	SPDZ	1	133	13056	0	-
	Conversion	2	271.73	63546	0	27030
	SPDZ-BMR	0	1.73	-	-	8383

Table 4. Two-party linear SVM: single-threaded (non-amortised) online phase costs and preprocessing costs with $\text{sec} = 64$.

Circuit type	Sub-Protocol	Preprocessing protocol (comm.)			Total
		LowGear	WRK (indep.)	WRK (dep.)	
SPDZ		54 MB	-	-	54 MB
GC		-	4917 MB	1768 MB	6685 MB
Marbled	SPDZ	24.48 MB	-	-	112.3 MB
	daBit convert	71.13 MB	9.43 MB	3.39 MB	
	GC	-	2.92 MB	1.05 MB	

Table 5. Two-party linear SVM communication cost for preprocessing in MBytes and statistical security $\text{sec} = 40$.

Acknowledgements

We would like to thank COED at COSIC at KU Leuven, and Emmanuela Orsini, Nigel Smart and Younes Talibi in particular, as well as Marcel Keller and Peter Scholl, for providing their enlightening insight and suggestions. Special thanks goes to Marcel Keller who helped with

the understanding of SPDZ-BMR code in the MP-SPDZ framework. We would also like to thank the anonymous reviewers for their helpful suggestions on how to improve the paper.

This work has been supported in part by ERC Advanced Grant ERC-2015-AdG-IMPACT, by the Defense Advanced Research Projects Agency (DARPA) and Space and Naval Warfare Systems Center, Pacific (SSC Pacific) under contract No. N66001-15-C-4070, and by the FWO under an Odysseus project GOH9718N.

References

- [ABF⁺18] Toshinori Araki, Assi Barak, Jun Furukawa, Marcel Keller, Yehuda Lindell, Kazuma Ohara, and Hikaru Tsuchida. Generalizing the spdz compiler for other protocols. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, pages 880–895. ACM, 2018.
- [AKO⁺18] A Aly, M Keller, E Orsini, D Rotaru, P Scholl, N Smart, and T Wood. Scale-mamba v1.3 : Documentation, 2018. <https://homes.esat.kuleuven.be/~nsmart/SCALE/>.
- [Ana19] N1 Analytics. MP-SPDZ, 2019. <https://github.com/n1analytics/MP-SPDZ>.
- [AOR⁺19] Abdelrahman Aly, Emmanuela Orsini, Dragos Rotaru, Nigel P. Smart, and Tim Wood. Zaphod: Efficiently combining lss and garbled circuits in scale. Cryptology ePrint Archive, Report 2019/974, 2019. <https://eprint.iacr.org/2019/974>.
- [BDK⁺18] Niklas Büscher, Daniel Demmler, Stefan Katzenbeisser, David Kretzmer, and Thomas Schneider. HyCC: Compilation of hybrid protocols for practical secure computation. In David Lie, Mohammad Mannan, Michael Backes, and XiaoFeng Wang, editors, *ACM CCS 2018*, pages 847–861. ACM Press, October 2018.
- [BDOZ11] Rikke Bendlin, Ivan Damgård, Claudio Orlandi, and Sarah Zakarias. Semi-homomorphic encryption and multiparty computation. In Kenneth G. Paterson, editor, *EUROCRYPT 2011*, volume 6632 of *LNCS*, pages 169–188. Springer, Heidelberg, May 2011.
- [Bea92] Donald Beaver. Efficient multiparty protocols using circuit randomization. In Joan Feigenbaum, editor, *CRYPTO’91*, volume 576 of *LNCS*, pages 420–432. Springer, Heidelberg, August 1992.
- [Ben18] Aner Ben-Efraim. On multiparty garbling of arithmetic circuits. In Thomas Peyrin and Steven Galbraith, editors, *ASIACRYPT 2018, Part III*, volume 11274 of *LNCS*, pages 3–33. Springer, Heidelberg, December 2018.
- [BGV11] Zvika Brakerski, Craig Gentry, and Vinod Vaikuntanathan. Fully homomorphic encryption without bootstrapping. Cryptology ePrint Archive, Report 2011/277, 2011. <http://eprint.iacr.org/2011/277>.
- [BMR90] Donald Beaver, Silvio Micali, and Phillip Rogaway. The round complexity of secure protocols (extended abstract). In *22nd ACM STOC*, pages 503–513. ACM Press, May 1990.

- [BMR16] Marshall Ball, Tal Malkin, and Mike Rosulek. Garbling gadgets for boolean and arithmetic circuits. In Edgar R. Weippl, Stefan Katzenbeisser, Christopher Kruegel, Andrew C. Myers, and Shai Halevi, editors, *ACM CCS 2016*, pages 565–577. ACM Press, October 2016.
- [Can00] Ran Canetti. Universally composable security: A new paradigm for cryptographic protocols. Cryptology ePrint Archive, Report 2000/067, 2000. <http://eprint.iacr.org/2000/067>.
- [Can01] Ran Canetti. Universally composable security: A new paradigm for cryptographic protocols. In *42nd FOCS*, pages 136–145. IEEE Computer Society Press, October 2001.
- [CDE⁺18] Ronald Cramer, Ivan Damgård, Daniel Escudero, Peter Scholl, and Chaoping Xing. SPD \mathbb{Z}_{2^k} : Efficient MPC mod 2^k for dishonest majority. In Hovav Shacham and Alexandra Boldyreva, editors, *CRYPTO 2018, Part II*, volume 10992 of *LNCS*, pages 769–798. Springer, Heidelberg, August 2018.
- [CKKZ12] Seung Geol Choi, Jonathan Katz, Ranjit Kumaresan, and Hong-Sheng Zhou. On the security of the “free-XOR” technique. In Ronald Cramer, editor, *TCC 2012*, volume 7194 of *LNCS*, pages 39–53. Springer, Heidelberg, March 2012.
- [DEF⁺19] Ivan Damgård, Daniel Escudero, Tore Frederiksen, Marcel Keller, Peter Scholl, and Nikolaj Volgushev. New primitives for actively-secure mpc over rings with applications to private machine learning. 2019.
- [DFK⁺06] Ivan Damgård, Matthias Fitzi, Eike Kiltz, Jesper Buus Nielsen, and Tomas Toft. Unconditionally secure constant-rounds multi-party computation for equality, comparison, bits and exponentiation. In *Theory of Cryptography Conference*, pages 285–304. Springer, 2006.
- [DKL⁺13] Ivan Damgård, Marcel Keller, Enrique Larraia, Valerio Pastro, Peter Scholl, and Nigel P Smart. Practical covertly secure mpc for dishonest majority—or: breaking the spdz limits. In *European Symposium on Research in Computer Security*, pages 1–18. Springer, 2013.
- [DPSZ12] Ivan Damgård, Valerio Pastro, Nigel Smart, and Sarah Zakarias. Multiparty computation from somewhat homomorphic encryption. In *Advances in Cryptology—CRYPTO 2012*, pages 643–662. Springer, 2012.
- [DSZ15] Daniel Demmler, Thomas Schneider, and Michael Zohner. ABY - A framework for efficient mixed-protocol secure two-party computation. In *NDSS 2015*. The Internet Society, February 2015.
- [DZ13] Ivan Damgård and Sarah Zakarias. Constant-overhead secure computation of boolean circuits using preprocessing. In *Theory of Cryptography*, pages 621–641. Springer, 2013.
- [FFFP04] Li Fei-Fei, R Fergus, and P Perona. Learning Generative Visual Models from Few Training Examples: An Incremental Bayesian Approach Tested on 101 Object Categories. In *CVPR*, pages 178–178. IEEE, 2004.
- [FKOS15] Tore Kasper Frederiksen, Marcel Keller, Emmanuela Orsini, and Peter Scholl. A unified approach to MPC with preprocessing using OT. In Tetsu Iwata and

- Jung Hee Cheon, editors, *ASIACRYPT 2015, Part I*, volume 9452 of *LNCS*, pages 711–735. Springer, Heidelberg, November / December 2015.
- [HKS⁺10] Wilko Henecka, Stefan Kögl, Ahmad-Reza Sadeghi, Thomas Schneider, and Immo Wehrenberg. TASTY: tool for automating secure two-party computations. In Ehab Al-Shaer, Angelos D. Keromytis, and Vitaly Shmatikov, editors, *ACM CCS 2010*, pages 451–462. ACM Press, October 2010.
 - [HOSS18] Carmit Hazay, Emmanuela Orsini, Peter Scholl, and Eduardo Soria-Vazquez. TinyKeys: A new approach to efficient multi-party computation. In Hovav Shacham and Alexandra Boldyreva, editors, *CRYPTO 2018, Part III*, volume 10993 of *LNCS*, pages 3–33. Springer, Heidelberg, August 2018.
 - [HSS17] Carmit Hazay, Peter Scholl, and Eduardo Soria-Vazquez. Low cost constant round MPC combining BMR and oblivious transfer. In Tsuyoshi Takagi and Thomas Peyrin, editors, *ASIACRYPT 2017, Part I*, volume 10624 of *LNCS*, pages 598–628. Springer, Heidelberg, December 2017.
 - [KOS16] Marcel Keller, Emmanuela Orsini, and Peter Scholl. MASCOT: faster malicious arithmetic secure computation with oblivious transfer. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, pages 830–842. ACM, 2016.
 - [KPR18] Marcel Keller, Valerio Pastro, and Dragos Rotaru. Overdrive: Making SPDZ great again. In Jesper Buus Nielsen and Vincent Rijmen, editors, *EUROCRYPT 2018, Part III*, volume 10822 of *LNCS*, pages 158–189. Springer, Heidelberg, April / May 2018.
 - [KS08] Vladimir Kolesnikov and Thomas Schneider. Improved garbled circuit: Free XOR gates and applications. In *International Colloquium on Automata, Languages, and Programming*, pages 486–498. Springer, 2008.
 - [KSS10] Vladimir Kolesnikov, Ahmad-Reza Sadeghi, and Thomas Schneider. From dust to dawn: Practically efficient two-party secure function evaluation protocols and their modular design. *IACR Cryptology ePrint Archive*, 2010:79, 2010.
 - [KSS13] Vladimir Kolesnikov, Ahmad-Reza Sadeghi, and Thomas Schneider. A systematic approach to practically efficient general two-party secure function evaluation protocols and their modular design. *Journal of Computer Security*, 21(2):283–315, 2013.
 - [KSS14] Florian Kerschbaum, Thomas Schneider, and Axel Schröpfer. Automatic protocol selection in secure two-party computations. In Ioana Boureanu, Philippe Owesarski, and Serge Vaudenay, editors, *ACNS 14*, volume 8479 of *LNCS*, pages 566–584. Springer, Heidelberg, June 2014.
 - [KY18] Marcel Keller and Avishay Yanai. Efficient maliciously secure multiparty computation for RAM. In Jesper Buus Nielsen and Vincent Rijmen, editors, *EUROCRYPT 2018, Part III*, volume 10822 of *LNCS*, pages 91–124. Springer, Heidelberg, April / May 2018.
 - [LPSY15] Yehuda Lindell, Benny Pinkas, Nigel P. Smart, and Avishay Yanai. Efficient constant round multi-party computation combining BMR and SPDZ. In Rosario

- Gennaro and Matthew J. B. Robshaw, editors, *CRYPTO 2015, Part II*, volume 9216 of *LNCS*, pages 319–338. Springer, Heidelberg, August 2015.
- [MR18] Payman Mohassel and Peter Rindal. Aby 3: a mixed protocol framework for machine learning. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, pages 35–52. ACM, 2018.
 - [MRSV19] Eleftheria Makri, Dragos Rotaru, Nigel P Smart, and Frederik Vercauteren. Epic: efficient private image classification (or: learning from the masters). In *Cryptographers’ Track at the RSA Conference*, pages 473–492. Springer, 2019.
 - [NNOB12] Jesper Buus Nielsen, Peter Sebastian Nordholt, Claudio Orlandi, and Sai She-shank Burra. A new approach to practical active-secure two-party computation. In *Advances in Cryptology–CRYPTO 2012*, pages 681–700. Springer, 2012.
 - [RWT⁺18] M. Sadegh Riazi, Christian Weinert, Oleksandr Tkachenko, Ebrahim M. Songhori, Thomas Schneider, and Farinaz Koushanfar. Chameleon: A hybrid secure computation framework for machine learning applications. In Jong Kim, Gail-Joon Ahn, Seungjoo Kim, Yongdae Kim, Javier López, and Taesoo Kim, editors, *ASIACCS 18*, pages 707–721. ACM Press, April 2018.
 - [SW19] Nigel P Smart and Tim Wood. Error detection in monotone span programs with application to communication-efficient multi-party computation. In *Cryptographers’ Track at the RSA Conference*, pages 210–229. Springer, 2019.
 - [WRK17] Xiao Wang, Samuel Ranellucci, and Jonathan Katz. Global-scale secure multi-party computation. In Bhavani M. Thuraisingham, David Evans, Tal Malkin, and Dongyan Xu, editors, *ACM CCS 2017*, pages 39–56. ACM Press, October / November 2017.
 - [Yao86] Andrew Chi-Chih Yao. How to generate and exchange secrets (extended abstract). In *27th FOCS*, pages 162–167. IEEE Computer Society Press, October 1986.

A Garbling

SPDZ-BMR Garbling Lindell et al. [LPSY15] gave generic multiparty method, known as SPDZ-BMR, for garbling in a constant number of rounds with malicious security where the preprocessed material is obtained from SPDZ [DPSZ12]. Their method is (roughly) to execute the classic [BMR90] multiparty garbling protocol using SPDZ to generate all the necessary secrets and to compute the ciphertexts. While the circuits are Boolean, the wires masks are arithmetic shares in \mathbb{F}_p of binary values and the wire keys random elements of \mathbb{F}_p secret-shared amongst the parties. Importantly, it was shown that it was not necessary for parties to provide zero-knowledge proofs that the evaluations of the pseudorandom function (PRF) used for encryption was done honestly, as the evaluators would abort with overwhelming probability in κ if parties cheated in this way.

The FreeXOR garbling technique [KS08] is an optimisation of Yao’s original garbling [Yao86, Oral presentation] that requires no data to be sent between the garbler and evaluator for an XOR gate, and crucially relies on the fact that the keys are elements of a field of characteristic 2. Towards the goal of a multiparty garbling protocol with FreeXOR, one

might hope to perform SPDZ-BMR over \mathbb{F}_{2^k} . One of the reasons this was not considered for SPDZ-BMR was (presumably) that the SPDZ offline phase was much faster for large prime fields than extension fields. Indeed, the most efficient variant of SPDZ used BGV [BGV11] as the SHE scheme, which meant that while the offline phase could parallelise through ciphertext packing, for large extension fields – and in particular for finite extensions of \mathbb{F}_2 – the amount of available packing was limited. However, shortly after this Keller et al. [KOS16] showed how to use oblivious transfer (OT) to perform the offline phase even more efficiently. This solution was shown to be more efficient than using SHE for extension fields. Despite recent work [KPR18] showing that SHE solutions outperform OT solutions for large prime fields, [KOS16] remains faster over extension fields. Subsequently, Keller and Yanai [KY18] showed how to apply FreeXOR in the multiparty setting using SPDZ-BMR-style garbling where the SPDZ shares are in \mathbb{F}_{2^k} instead of \mathbb{F}_p .

Meanwhile, Hazay et al. [HSS17] also showed how to obtain FreeXOR in the multiparty setting, again over \mathbb{F}_{2^k} , but take a different approach from SPDZ-BMR: they do not make use of a full-blown MPC functionality and instead produce an *unauthenticated* garbled circuit – it is merely additively shared, whereas in SPDZ-BMR and [KY18], the garbled circuit is authenticated with MACs. Active security comes from the fact that an incorrectly-garbled circuit will only cause the parties to abort when evaluating it. This approach requires only a single (authenticated) \mathbb{F}_2 multiplication per AND gate.

We use the multiparty Boolean circuit garbling protocol [KY18] for our implementation. The [KY18] protocol is less efficient than [HSS17] and [WRK17], but the implementation [Ana19] is easier to integrate with the SPDZ compiler to be able to switch between different online phases of an MPC program [ABF⁺18]. Despite this we have good reason to believe that the generation of the specialised preprocessing required in our solution dovetails with most if not all of these alternative garbling schemes as the only requirements are the following:

- Parties should be able to authenticate their own secret inputs (in fact, secret bits suffices), for whatever authentication method is used in the protocol.
- Parties should be able to compute the XOR of authenticated bits.

Unfortunately, the authentication is usually abstracted away garbling functionalities so we cannot make straightforward claims about using garbling in a black-box way.

Encryption Before we describe the garbling, we first briefly describe the encryption scheme used in the protocol. The formalism of the security of using FreeXOR in the multiparty context was not given in [KY18] so we provide an overview here based on [HSS17]. In the garbling protocol, messages are encrypted by computing the XOR of the message with a pseudorandom one-time-pad generated by a PRF under keys held by multiple parties. The SPDZ-BMR technique of encryption requires a PRF which is a *pseudorandom function under multiple keys* [LPSY15, Defn 1] (see Definition D.1 in Appendix D). In order to use the FreeXOR technique, a stronger assumption is needed: *circular 2-correlation robust*. Detailed analysis of a similar assumption for hash functions was given by [CKKZ12], and Hazay et al. [HSS17] gave a variant for PRFs, of which we now give an outline.

To make the definition, we define the following oracles. Sample $R \xleftarrow{\$} \{0, 1\}^\kappa$ and let $|C|$ denote the number of gates in the circuit. For a PRF $F : \{0, 1\}^\kappa \times \{0, 1\}^\kappa \times \{0, 1\}^{\log |C| + \log n} \rightarrow \{0, 1\}^\kappa$, define two oracles which take inputs from the set $\{0, 1\}^\kappa \times \{0, 1\}^\kappa \times [|C|] \times [n] \times \{0, 1\}^3$:

- Oracle Circ_R : on input $(k_1, k_2, g, j, b_1, b_2, b_3)$, return $F_{k_1 \oplus b_1 \cdot R, k_2 \oplus b_2 \cdot R}(g \| j) \oplus b_3 \cdot R$.
- Oracle RO : on input $(k_1, k_2, g, j, b_1, b_2, b_3)$, if the message has been queried before, output whatever was given last time; otherwise, sample a random string $k_3 \xleftarrow{\$} \{0, 1\}^\kappa$ and output k_3 .

Definition A.1 (See [HSS17, Defn 2.3]). *A PRF F is called circular 2-correlation robust if for any non-uniform polynomial-time distinguisher \mathcal{D} making legal queries (see below) to its oracle, there exists a negligible function ν such that*

$$\left| \Pr[R \xleftarrow{\$} \{0, 1\}^\kappa; \mathcal{D}^{\text{Circ}_R(\cdot)}(1^\kappa) = 1] - \Pr[\mathcal{D}^{\text{RO}(\cdot)}(1^\kappa) = 1] \right| \leq \nu(\kappa).$$

“Legal queries” are defined as any query except the following:

- $(k_1, k_2, g, j, 0, 0, b_3)$ (otherwise the distinguisher can distinguish trivially as it learns $F_{k_1, k_2}(g \| j)$ and it can compute this on its own).
- The query $(k_1, k_2, g, j, b_1, b_2, 1 - b_3)$ if the query $(k_1, k_2, g, j, b_1, b_2, b_3)$ has already been made (otherwise the global difference is leaked).

The encryption of a message $m \in \{0, 1\}^\kappa$ under keys $k_u \leftarrow k_u^1 \parallel \dots \parallel k_u^n$ and $k_v \leftarrow k_v^1 \parallel \dots \parallel k_v^n$ and nonce r , where P_i holds they keys $k_u^i, k_v^i \in \{0, 1\}^\kappa$, is defined as

$$\text{Enc}_{k_u, k_v}(m; r) \leftarrow \left(\bigoplus_{j=1}^n F_{k_u^j, k_v^j}(r) \right) \oplus m.$$

For clarity, we will write the formula explicitly in the circuit garbling rather than abstract to the encryption notation.

Garbling Protocol The high-level view is as follows. Let $g : \{0, 1\}^2 \rightarrow \{0, 1\}$ denote the gate g . In circuit garbling, first the garbler samples a “zero key” and a “one key” for every wire in the circuit. A wire connection exiting one gate and entering another is considered one wire, as are all circuit input and output wires. Then, each Boolean fan-in-two gate with input wires u, v and output wire w is converted to a set of 4 ciphertexts in the following way: for each $(\alpha, \beta) \in \{0, 1\}^2$ the garbler encrypts the key $k_{w, g(\alpha, \beta)}$ under the pair of input keys $k_{u, \alpha}$ and $k_{v, \beta}$. The garbler converts all gates to quadruples of ciphertexts. Note that this can be done in parallel for all gates. To evaluate the circuit, the evaluator is given the circuit input wire keys corresponding to his inputs (obviously), and using these can decrypt one ciphertext in each quadruple (gate) to obtain the final output. Other measures, as we describe below, are required to hide information about intermediate wire values as the circuit is being garbled, but this is the basic outline.

In BMR garbling, every party acts as both garbler and evaluator. Each party generates a circuit for which it knows all the wire keys, but where each ciphertext is encrypted under *all* parties’ corresponding wire keys. This means that every party must evaluate all n circuits

in parallel to decrypt each subsequent gate and so learn all n succeeding keys. The idea of SPDZ-BMR garbling is to use MPC to compute the ciphertexts. The full protocol, modified for our purposes and incorporating FreeXOR, is provided in Figures 16, 17, 18 and 19, but we give an overview of correctness here. The functionality \mathcal{F}_{MPC} is part of $\mathcal{F}_{\text{Prep}}$ in Figure 3. The parties first call \mathcal{F}_{MPC} with input $(\text{Initialise}, \mathbb{F}_{2^k}, 1)$ where 1 is a session identifier.

Global difference Once at the beginning of the protocol, for each $i \in [n]$ the parties call \mathcal{F}_{MPC} with input $(\text{RElt}, \llbracket R^i \rrbracket_{2^k}, 1)$ and then $(\text{Open}, i, \llbracket R^i \rrbracket_{2^k}, 1)$ so that P_i obtains a random $R^i \in \mathbb{F}_{2^k}$, called its “global difference”. These are used later for defining keys in a special way.

Wire Masks For each wire in the circuit, the parties call \mathcal{F}_{MPC} with input $(\text{RBit}, \llbracket \lambda_w \rrbracket_{2^k}, 1)$. These are known as the *masking* or *permutation* bits and are used to permute the four ciphertexts in each gate, which is necessary to hide intermediate wire values (which leak information on the circuit inputs). Instead of evaluating $g(\alpha, \beta)$ on the actual inputs ρ and σ , evaluators hold “signal bits” $\Lambda_u \leftarrow \rho \oplus \lambda_u$ and $\Lambda_v \leftarrow \sigma \oplus \lambda_v$ and compute $\Lambda_w \leftarrow \tilde{g}(\Lambda_u, \Lambda_v)$ where \tilde{g} is defined as $(\alpha, \beta) \mapsto g(\alpha \oplus \lambda_u, \beta \oplus \lambda_v) \oplus \lambda_w$ where λ_w is the mask for the output wire⁴.

Garbling

- For an AND gate g with input wires u and v and output wire w , for each $i \in [n]$ the parties call \mathcal{F}_{MPC} with input $(\text{RElt}, \llbracket k_{u,0}^i \rrbracket_{2^k}, 1)$ to obtain a random field element and open it to party P_i . The one key $k_{u,1}^i$ is set to be $k_{u,0}^i \oplus R^i$ by P_i and the parties compute $\llbracket k_{u,1}^i \rrbracket_{2^k} \leftarrow \llbracket k_{u,0}^i \rrbracket_{2^k} \oplus \llbracket R^i \rrbracket_{2^k}$. The parties do the same for v and w . It is possible to garble with random one keys as described in the overview above, but this global difference allows more efficient garbling and evaluation as outlined below with no loss to security. Then each party evaluates the PRF on the four combinations of their own input keys $\{(k_{u,\alpha}^i, k_{v,\beta}^i) : \alpha, \beta \in \{0, 1\}\}$ and calls \mathcal{F}_{MPC} with input

$$(\text{Input}, i, \llbracket F_{\alpha,\beta}^{g,i,j} \rrbracket_{2^k}, F_{k_{u,\alpha}^i, k_{v,\beta}^i}(g\|j), 1)$$

which form $4n$ authenticated (pseudorandom one-time-pad) encryption keys, indexed by $\alpha, \beta \in \{0, 1\}$ and $j \in [n]$. In MPC, the parties then compute, for all $(\alpha, \beta) \in \{0, 1\}^2$ and all $j \in [n]$, the ciphertext

$$\llbracket \tilde{g}_{\alpha,\beta}^j \rrbracket_{2^k} \leftarrow \left(\bigoplus_{i=1}^n \llbracket F_{\alpha,\beta}^{g,i,j} \rrbracket_{2^k} \right) \oplus \llbracket k_{w,0}^j \rrbracket_{2^k} \oplus \llbracket R^j \rrbracket_{2^k} \cdot ((\llbracket \lambda_u \rrbracket_{2^k} \oplus \alpha) \cdot (\llbracket \lambda_v \rrbracket_{2^k} \oplus \beta) \oplus \llbracket \lambda_w \rrbracket_{2^k}).$$

In other words, for each $j \in [n]$ the parties compute an encryption in MPC of wire key $k_{w,0}^j \oplus R^j \cdot ((\lambda_u \oplus \alpha) \cdot (\lambda_v \oplus \beta) \oplus \lambda_w)$ under all four possible pairs of input keys of every other party. Note that the j^{th} set of four ciphertexts are permuted by the *same* masks for every j .

⁴ This is equivalent to randomly permuting the four ciphertexts (indexed by $\{1, 2, 3, 4\}$) by the (secret) permutation $((13)(24))^{\lambda_u}((12)(34))^{\lambda_v}$.

- For an XOR gate, the parties set $k_{w,0}^j \leftarrow k_{u,0}^j \oplus k_{v,0}^j$ for all j and the output mask as $\lambda_w \leftarrow \lambda_u \oplus \lambda_v$.

After all the garbling is performed, all the ciphertexts (currently held in the MPC engine) are opened.

Input For a party P_i to provide an input $x \in \{0, 1\}$ on wire w , the parties open $\llbracket \lambda_w \rrbracket_{2^k}$ to P_i and then P_i broadcasts $\Lambda_w \leftarrow x \oplus \lambda_w$, known as a *signal bit*. For all $j \in [n]$, P_j broadcasts k_{w,Λ_w}^j .

Evaluation After the n keys and signal bits, one for each input wire, are obtained, the parties do the following.

- For an AND gate, for every $j \in [n]$, each party computes the n succeeding wire keys as

$$\begin{aligned} k_{w,\cdot}^j &\leftarrow \tilde{g}_{\Lambda_u, \Lambda_v}^j \oplus \bigoplus_{i=1}^n F_{k_{u,\Lambda_u}^i, k_{v,\Lambda_v}^i}(g \parallel j) \\ &= k_{w,0}^j \oplus R^j \cdot ((\lambda_u \oplus \Lambda_u) \cdot (\lambda_v \oplus \Lambda_v) \oplus \lambda_w) \end{aligned}$$

Then each P_i compares $k_{w,\cdot}^i$ to the two keys $k_{w,0}^i$ and $k_{w,1}^i$ that it owns in order to determine the signal bit Λ_w . By the security of the PRF (as argued in SPDZ-BMR), the ciphertexts corresponding to the other keys cannot be decrypted. Observe that since $\Lambda_u = x \oplus \lambda_u$ and $\Lambda_v = y \oplus \lambda_v$ where x and y are the actual inputs, the resulting keys (for $j \in [n]$) are $k_{w,0}^j \oplus R^j \cdot (x \cdot y \oplus \lambda_w)$, which are exactly k_{w,Λ_w}^j where $\Lambda_w = x \cdot y \oplus \lambda_w$.

- For an XOR gate, every party computes the n output signal bit as $\Lambda_w \leftarrow \Lambda_u \oplus \Lambda_v$ and sets the keys as $k_{w,\Lambda_w}^j \leftarrow k_{u,\Lambda_u}^j \oplus k_{v,\Lambda_v}^j$. Observe that

$$\Lambda_w \leftarrow \Lambda_u \oplus \Lambda_v = (x \oplus \lambda_u) \oplus (y \oplus \lambda_v) = (x \oplus y) \oplus (\lambda_u \oplus \lambda_v) = (x \oplus y) \oplus \lambda_w$$

by the way the masks are constructed; similarly,

$$k_{w,\Lambda_w}^j \leftarrow k_{u,\Lambda_u}^j \oplus k_{v,\Lambda_v}^j = (k_{u,0}^j \oplus R^j \cdot \Lambda_u) \oplus (k_{v,0}^j \oplus R^j \cdot \Lambda_v) = k_{w,0}^j \oplus R^j \cdot (\Lambda_u \oplus \Lambda_v) = k_{w,\Lambda_w}^j.$$

Output In the usual BMR protocol, immediately after garbling, the parties open the masks for output wires. This enables all parties to view the output bits. We will assume the parties simply hold the final output in secret-shared form, which they can do simply by not opening the output masks, and by computing (locally) the XOR of the public signal bit with the output mask. I.e. an output bit is shared as

$$\llbracket b \rrbracket_{2^k} \leftarrow \llbracket \lambda_w \rrbracket_{2^k} \oplus \Lambda_w.$$

In Section 4.2 we describe the modifications necessary to this standard garbling technique to provide inputs from get outputs to \mathbb{F}_p .

B Various Functionalities and Protocols

In Figures 11, 12, 13 and 14 we give some of the standard functionalities and protocols securely realising them in the UC framework. We omit the proofs as they are standard.

Protocol Π_{Rand}

This protocol is in the $\mathcal{F}_{\text{Commit}}$ -hybrid model. Let $\text{RShuffle}(\text{seed}, s)$ denote any deterministic algorithm that takes a random seed seed and a vector s and outputs a permutation of components of s . Recall κ is the computational security parameter.

Initialise Parties agree on a session identifier sid and call $\mathcal{F}_{\text{Commit}}$ with input $(\text{Initialise}, \{0, 1\}^\kappa, \text{sid})$.

Seed This is a subroutine. Parties do the following:

1. For each $i \in [n]$, P_i samples $\text{seed}_i \xleftarrow{\$} \{0, 1\}^\kappa$.
2. For each $i \in [n]$, P_i sends the message $(\text{Commit}, i, \text{seed}_i, \text{sid})$ to $\mathcal{F}_{\text{Commit}}$ and all other parties send $(\text{Commit}, i, \perp, \text{sid})$; all parties receive $\text{id}_{\text{seed}_i}$ in response.
3. All parties call $(\text{Open}, i, \text{id}_{\text{seed}_i}, \text{sid})$ to obtain $\{\text{seed}_i\}_{i \in [n]}$.
4. P_i sets $\text{seed} \leftarrow \bigoplus_{i=1}^n \text{seed}_i$.

Random subset To compute a random subset of size t of a set X , parties run **Seed** to obtain a seed seed for a PRG and then do the following:

1. Let $X = \{x_i\}_{i=1}^{|X|}$. Parties set the vector $s = (s_1, \dots, s_{|X|}) \leftarrow (1, \dots, 1, 0, \dots, 0) \in \{0, 1\}^{|X|}$, where the first t bits are set to 1 and the remaining bits set to 0.
2. Each P_i locally computes $s' = (s'_1, \dots, s'_{|X|}) \leftarrow \text{RShuffle}(\text{seed}, s)$ and outputs the set $S \leftarrow \{x_i : s'_i = 1\}$.

Random buckets To put a set of items indexed by a set X into buckets of size t where t divides $|X|$, parties run **Seed** to obtain a seed seed for a PRG and then do the following:

1. Let $X = \{x_i\}_{i=1}^{|X|}$. Each P_i locally computes $s' \leftarrow \text{RShuffle}(\text{seed}, s)$ where $s \leftarrow (i)_{i=1}^{|S|}$.
2. For each $i = 1$ to $|S|/t$, let $S_i \leftarrow \{x_{s'_j} : (i-1) \cdot t < j \leq i \cdot t\}$.

Fig. 11. Protocol Π_{Rand}

Functionality $\mathcal{F}_{\text{Commit}}$

The functionality keeps track of the current session using a session identifier **sid**. If a party provides an input with **sid** different from what was sent in **Initialise**, the functionality outputs **Reject** to all parties and awaits another message. Recall that κ is the computational security parameter.

Initialise On input (**Initialise**, X , **sid**) from all parties where X is a set, initialise a dictionary of values **Val** with identifiers **Val.Keys**.

Commit On input (**Commit**, i , x , **sid**) from P_i , or the adversary if P_i is corrupt, and (**Commit**, i , \perp , **sid**) from all other parties, where $x \in X$, choose new identifier $\text{id}_x \xleftarrow{\$} \{0, 1\}^\kappa$, add id_x to **Val.Keys**, set $\text{Val}[\text{id}_x] \leftarrow x$, and send id_x to all parties.

Open On input (**Open**, i , id_x , **sid**) from all parties where $\text{id}_x \in \text{Val.Keys}$, if P_i is corrupt then await a message **OK** or **Reject** from the adversary. If the message is **OK** or P_i is honest then send $\text{Val}[\text{id}_x]$ to all parties and otherwise halt.

Fig. 12. Functionality $\mathcal{F}_{\text{Commit}}$

Protocol Π_{Commit}

This protocol is in the \mathcal{F}_{RO} -hybrid model (ROM). Recall that κ is the computational security parameter. We write $a\|b$ to mean a concatenated with b .

Initialise The parties agree on a session identifier **sid** and call \mathcal{F}_{RO} with input (**Initialise**, $\{0, 1\}^\kappa$, **sid**).

Commit For P_i to commit to x , the parties do the following:

1. P_i samples $r \xleftarrow{\$} \{0, 1\}^\kappa$
2. P_i calls \mathcal{F}_{RO} with input (**Query**, $x\|r$, **sid**), receives h_x in response and broadcasts it.
3. All other parties store h_x locally.

Open For P_i to open the commitment to x , the parties do the following:

1. P_i broadcasts x and r .
2. P_j calls \mathcal{F}_{RO} with input (**Query**, $x\|r$, **sid**) and checks that the response is equal to h_x received during commitment.

Fig. 13. Protocol Π_{Commit}

Functionality \mathcal{F}_{RO}

The functionality keeps track of the current session using a session identifier **sid**. If a party provides an input with **sid** different from what was sent in **Initialise**, the functionality outputs **Reject** to all parties and awaits another message.

Initialise On input (**Initialise**, X , **sid**) from all parties, initialise a dictionary of values **Val**.

Random Element On input (**Query**, q , **sid**) from party P_i where $q \in \{0, 1\}^*$, if **Val**[q] has not yet been defined, uniformly sample $\text{Val}[q] \xleftarrow{\$} X$ and send **Val**[q] to P_i .

Fig. 14. Functionality \mathcal{F}_{RO}

C Switching Protocol for SPDZ-BMR

Theorem C.1. $\Pi_{\text{ABB+BMR}}$ securely realises $\mathcal{F}_{\text{CABB}}$ in the $\mathcal{F}_{\text{Prep}}$ -hybrid model.

Proof. We define a simulator in Figure 15. Note that the only secrets to which the simulator is not privy are the secret inputs of honest parties. Everything else in the protocol involves calls to $\mathcal{F}_{\text{Prep}}$ which is locally emulated by the simulator.

We define the hybrids in the same way as in the proof of Theorem 4.1. Suppose the adversary corrupts $t < n$ parties in total, indexed by a set A . We define a sequence of hybrid worlds (**Hybrid** h) $_{h=0}^{n-t}$ and show that each is indistinguishable from the previous. **Hybrid** h is defined as:

Hybrid h The simulator has the actual input of $n - t - h$ honest parties and must simulate the remaining h honest parties towards the adversary.

The simulator is described in Figure 15, parameterised by h .

Simulator $\mathcal{S}_{\text{ABB+BMR}}$

Let H_R denote the indexing set of the h honest parties whose inputs are known to $\mathcal{S}_{\text{ABB+BMR}}^h$. Initialise and run an internal copy of $\mathcal{F}_{\text{Prep}}$ with the adversary, answering every query by executing the code of $\mathcal{F}_{\text{Prep}}$. For simplicity of relaying messages between \mathcal{A} and $\mathcal{F}_{\text{CABB}}$, we assume $\text{sid} = 0$ for $\mathcal{F}_{\text{CABB}}$ as well as the \mathbb{F}_p instance of $\mathcal{F}_{\text{Prep}}$.

Initialise Initialise a local copy of $\mathcal{F}_{\text{Prep}}$ and await the inputs (**Initialise**, \mathbb{F}_p , 0) and (**Initialise**, \mathbb{F}_{2^k} , 1) from \mathcal{A} .

ABB For calls to $\mathcal{F}_{\text{Prep}}$ with $\text{sid} = 0$:

Input On input (**Input**, i , id , 0), forward the message to $\mathcal{F}_{\text{CABB}}$.

Add On input (**Add**, id_x , id_y , id , 0), forward the message to $\mathcal{F}_{\text{CABB}}$.

Multiply On input (**Multiply**, id_x , id_y , id , 0), forward the message to $\mathcal{F}_{\text{CABB}}$.

Output 1. Await a message (**Check**, 0) from \mathcal{A} to $\mathcal{F}_{\text{Prep}}$ and then await a message (**OK**, 0) or (**Abort**, 0) from \mathcal{A} , and if it is (**Abort**, 0) then ignore all further calls to $\mathcal{F}_{\text{Prep}}$ with $\text{sid} = 0$ and otherwise continue.

2. On input (**Open**, 0, id , 0) to $\mathcal{F}_{\text{Prep}}$, send the message (**Output**, id , 0) to $\mathcal{F}_{\text{CABB}}$ and relay the response x to \mathcal{A} .

3. Await a reply $x + \varepsilon$ from \mathcal{A} and the call by \mathcal{A} to $\mathcal{F}_{\text{Prep}}$ with input (**Check**, 0).

4. Await another message (**OK**, 0) or (**Abort**, 0) from \mathcal{A} . If $\varepsilon = 0$ and the message was (**OK**, 0) then send **OK** to $\mathcal{F}_{\text{CABB}}$, and otherwise send **Abort** to $\mathcal{F}_{\text{CABB}}$ and (**Abort**, 0) to \mathcal{A} and ignore all further calls to $\mathcal{F}_{\text{Prep}}$ with $\text{sid} = 0$ and otherwise continue.

Initialise garbling Run **Initialise** from $\Pi_{\text{ABB+BMR}}$ with \mathcal{A} .

Input layer Run **Input layer** from $\Pi_{\text{ABB+BMR}}$ with \mathcal{A} .

Garble Run $\Pi_{\text{ABB+BMR}}^{\text{Garble}}$ with \mathcal{A} .

Output layer Run **Output layer** from $\Pi_{\text{ABB+BMR}}$ with \mathcal{A} .

Open Run **Open** from $\Pi_{\text{ABB+BMR}}$ with \mathcal{A} .

Evaluate Suppose a circuit input x is from an honest party's input.

1. Send the message (**EvaluateCircuit**, C , $\text{id}_1, \dots, \text{id}_t$, id , 0) to $\mathcal{F}_{\text{CABB}}$.
2. Await the call (**Open**, 0, $\llbracket a - r \rrbracket$) from \mathcal{A} . Then:
 - If a is an input dependent on honest parties' inputs known to $\mathcal{S}_{\text{ABB+BMR}}^h$, retrieve the bits of the mask r from memory and compute and send $x \leftarrow a - r$ to the adversary.
 - If a is dependent on one or more honest parties' inputs then sample $x \leftarrow r' \xleftarrow{\$} \mathbb{F}_p$ and send it to \mathcal{A} .
3. Await a response $x + \varepsilon$ and if $\varepsilon \neq 0$ then send (**Abort**, 0) to \mathcal{A} and **Abort** to $\mathcal{F}_{\text{CABB}}$ and terminate; otherwise, continue.
4. Await the calls to $\mathcal{F}_{\text{Prep}}$ with input (**Open**, 0, $\llbracket \mathbf{k}_{u_j, x_j} \rrbracket_{2^k}, 1 \rangle_{j=0}^{\lceil \log p \rceil - 1}$ from \mathcal{A} , where u_j is the wire for the j^{th} input bit x_j of x , and respond honestly.
5. The simulator computes what honest parties would compute in the circuit evaluation. If an honest party would have aborted then the simulator sends **Abort** to $\mathcal{F}_{\text{CABB}}$, and otherwise sends **OK** and continues.

Fig. 15. Simulator $\mathcal{S}_{\text{ABB+BMR}}^h$

Claim. The $\mathcal{F}_{\text{Pprep}}$ -hybrid world is indistinguishable from **Hybrid** 0.

Proof. For the emulation of $\mathcal{F}_{\text{Pprep}}$ with $\text{sid} = 0$, the simulation is perfect.

For the emulation of $\mathcal{F}_{\text{Pprep}}$ with $\text{sid} = 1$, there are no private inputs of honest parties to the garbling, so it only remains to show that the transcript during evaluation reveals nothing about the (honest) parties' inputs.

In this hybrid, the simulator has access to all honest parties' inputs, so the simulator follows the protocol exactly in **Evaluate**, so the worlds are indistinguishable. ■

Claim. **Hybrid** h is indistinguishable from **Hybrid** $h + 1$ for $h = 0, \dots, n - t - 1$.

Proof. For the emulation of $\mathcal{F}_{\text{Pprep}}$ with $\text{sid} = 0$ and for the garbling the simulation is still perfect since honest parties' inputs are not required.

Indeed, the only call to **Input** is when the parties call $\mathcal{F}_{\text{Pprep}}$ during **Garble** for the PRF evaluations. The simulator can perform the PRF evaluations locally since the keys and global differences for honest parties are obtained from the emulation of $\mathcal{F}_{\text{Pprep}}$.

The only part of the circuit evaluation that may depend on honest parties' inputs is in **Evaluate**. Since the secret masks $\llbracket r \rrbracket_p$ are constructed from uniformly-sampled bits, by Lemma 2.1 the distribution of the uniformly sample $r' \xleftarrow{\$} \mathbb{F}_p$ is statistically close to the distribution of $a - r \pmod p$ where a is the input of an honest party and $r \xleftarrow{\$} \{0, 1\}^{\lceil \log p \rceil}$.

Now since the masking bits are sampled uniformly during the garbling and are unknown to the adversary (or environment), and the circuit can only be evaluated once, the intermediate wire values and the final bit (masked) outputs reveal nothing about the initial inputs to the circuit. Indeed, after evaluating the circuit the parties just have an identifier corresponding to the circuit output, which reveals no information on the underlying value by definition of the functionality. Thus the environment cannot use the evaluation to distinguish between the circuit evaluated on the actual value $a - r$ and the circuit evaluated on the sampled value random r' , even though the parties will obtain an “incorrect” output identifier with high probability. The simulator easily deals with this by obtaining the output from $\mathcal{F}_{\text{CABB}}$ and sending this to \mathcal{A} , ensuring that if the evaluation of the garbled circuit did caused an honest party to abort then the interaction with \mathcal{A} and $\mathcal{F}_{\text{CABB}}$ also abort. ■

Since $\mathcal{F}_{\text{Pprep}}$ is secure for $t = n - 1$, the result follows. □

Protocol $\Pi_{\text{ABB+BMR}}$

This protocol is secure in the $\mathcal{F}_{\text{Prep}}$ -hybrid model.

Initialise The parties call $\mathcal{F}_{\text{Prep}}$ with inputs $(\text{Initialise}, \mathbb{F}_p, 0)$ and $(\text{Initialise}, \mathbb{F}_{2^k}, 1)$.

Arithmetic

Input For P_i to provide input $x \in \mathbb{F}_p$, P_i calls $\mathcal{F}_{\text{Prep}}$ with input $(\text{Input}, i, \llbracket x \rrbracket_p, x, 0)$ and all other parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Input}, P_i, \llbracket x \rrbracket_p, \perp, 0)$, where $\llbracket x \rrbracket_p$ is a fresh identifier.

Add To add secrets x and y , parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Add}, \llbracket x \rrbracket_p, \llbracket y \rrbracket_p, \llbracket z \rrbracket_p, 0)$ where $\llbracket z \rrbracket_p$ is a new identifier.

Multiply To multiply secrets x and y , parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Add}, \llbracket x \rrbracket_p, \llbracket y \rrbracket_p, \llbracket z \rrbracket_p, 0)$ where $\llbracket z \rrbracket_p$ is a new identifier.

Output To receive output with identifier id , parties do the following:

1. The parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Check}, 0)$.
2. The parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, 0, \text{id}, 0)$.
3. The parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Check}, 0)$.

Circuit (All of the following procedures are performed, in order.)

Initialise garbling To garble a Boolean circuit C with identifiers W for wires, G_{AND} for AND gates and G_{XOR} for XOR gates, the parties do the following:

1. The parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{daBits}, \{\llbracket \lambda_w \rrbracket_{p,2^k}\}_{w \in W_o}, 0, 1)$ where W_o denotes the set indexing circuit output wires.
2. For each $i \in [n]$, the parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{RElt}, \llbracket R^i \rrbracket_{2^k}, 1)$ and then call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, i, \llbracket R^i \rrbracket_{2^k}, 1)$ to reveal R^i to P_i .

Input layer Let the number of \mathbb{F}_p inputs to the circuit be t . The parties do the following:

1. Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{daBits}, (\{\llbracket r_{i,j} \rrbracket_{p,2^k}\}_{j=0}^{\lceil \log p \rceil - 1})_{i=1}^t, 0, 1)$.
2. Set $\llbracket r_i \rrbracket_p \leftarrow \sum_{j=0}^{\lceil \log p \rceil - 1} 2^j \llbracket r_{i,j} \rrbracket_p$.
3. For $i = 1, \dots, t$, create the circuit $\text{ADDMOD}(x_i, y_i, p)$ and prepend these circuits to the circuit C to be garbled, augmenting G_{AND} and G_{XOR} as appropriate. See Section 4.2 for details.
4. For each input wire $w \in W$, for each $i \in [n]$,
 - (a) Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{RElt}, \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k}, 1)$.
 - (b) Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, i, \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k}, 1)$ to reveal $\mathbf{k}_{w,0}^i$ to P_i .
 - (c) P_i sets the one key as $\mathbf{k}_{w,1}^i \leftarrow \mathbf{k}_{w,0}^i \oplus R^i$ and the parties set $\llbracket \mathbf{k}_{w,1}^i \rrbracket_{2^k} \leftarrow \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k} \oplus \llbracket R^i \rrbracket_{2^k}$.
5. For every input wire w corresponding to an input $x_{i,j}$ of $\text{ADDMOD}(x_i, y_i, p)$, set $\lambda_{w,i,j} \leftarrow 0$.
6. For every input wire w corresponding to an input $y_{i,j}$ of $\text{ADDMOD}(x_i, y_i, p)$, call $\mathcal{F}_{\text{Prep}}$ with input $(\text{RBit}, \llbracket \lambda_w \rrbracket_{2^k}, 1)$ followed by $(\text{Open}, 0, \llbracket r_{i,j} \rrbracket_{2^k} \oplus \llbracket \lambda_{w,i,j} \rrbracket_{2^k}, 1)$ and store this as $\lambda_{w,i,j}$. Then for every $l \in [n]$, P_l sends $k_{w,i,j,\lambda_{w,i,j}}^l$ to all other parties.

Garble Refer to $\Pi_{\text{ABB+BMR}}^{\text{Garble}}$ in Figure 18.

(continued...)

40
Fig. 16. Protocol $\Pi_{\text{ABB+BMR}}$

Protocol $\Pi_{\text{ABB+BMR}}$ (continued)

Output layer For every wire w that is an (external, circuit) output wire, the parties do the following

1. Retrieve a daBit $\llbracket \lambda_{w'} \rrbracket_{p,2^k}$ from memory, generated in **Initialise**.
2. Compute $\llbracket \lambda_{w_0} \rrbracket_{2^k} \leftarrow \llbracket \lambda_w \rrbracket_{2^k} \oplus \llbracket \lambda_{w'} \rrbracket_{2^k}$.
3. Call $\mathcal{F}_{\text{Prep}}$ with input (Open, 0, $\llbracket \lambda_{w_0} \rrbracket_{2^k}, 1$); all parties store this locally in memory as the value A_{w_0} .

Open To open the circuit, the parties do the following:

1. For all $i \in [n]$, call $\mathcal{F}_{\text{Prep}}$ with input (Open, 0, $\llbracket \tilde{g}_{\alpha,\beta}^j \rrbracket_{2^k}, 1$) for all $g \in G_{\text{AND}}$, for all $j \in [n]$, for all $(\alpha, \beta) \in \{0, 1\}^2$. If the functionality returns \perp , the parties abort, and otherwise the parties (locally) output $((\tilde{g}_{0,0}^j, \tilde{g}_{0,1}^j, \tilde{g}_{1,0}^j, \tilde{g}_{1,1}^j)_{j=1}^n)_{g \in G}$ and the input mask identifiers $\llbracket r_1 \rrbracket_p, \dots, \llbracket r_t \rrbracket_p$.
2. Call $\mathcal{F}_{\text{Prep}}$ with input (Check, 1).

Evaluate Refer to $\Pi_{\text{ABB+BMR}}^{\text{Eval}}$ in Figure 19.

Fig. 17. Protocol $\Pi_{\text{ABB+BMR}}$ (continued)

Subprotocol $\Pi_{\text{ABB+BMR}}^{\text{Garble}}$

Garble Traversing the circuit in topological order, for every gate $g \in G$ with (internal) input wires u and v and (internal) output wire w ,

- If g is an XOR gate, i.e. $g \in G_{\text{XOR}}$,
 1. The parties set $\llbracket \lambda_w \rrbracket_{2^k} \leftarrow \llbracket \lambda_u \rrbracket_{2^k} \oplus \llbracket \lambda_v \rrbracket_{2^k}$.
 2. For each $i \in [n]$, P_i computes $\mathbf{k}_{w,0}^i \leftarrow \mathbf{k}_{u,0}^i \oplus \mathbf{k}_{v,0}^i$ and $\mathbf{k}_{w,1}^i \leftarrow \mathbf{k}_{w,0}^i \oplus R^i$ and all parties set $\llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k} \leftarrow \llbracket \mathbf{k}_{u,0}^i \rrbracket_{2^k} \oplus \llbracket \mathbf{k}_{v,0}^i \rrbracket_{2^k}$ and $\llbracket \mathbf{k}_{w,1}^i \rrbracket_{2^k} \leftarrow \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k} \oplus \llbracket R^i \rrbracket_{2^k}$.
- If g is an AND gate, i.e. $g \in G_{\text{AND}}$,
 1. The parties call $\mathcal{F}_{\text{Prep}}$ with input $(\text{RBit}, \llbracket \lambda_w \rrbracket_{2^k}, 1)$.
 2. For each $i \in [n]$,
 - (a) Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{RElt}, \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k}, 1)$.
 - (b) Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, i, \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k}, 1)$ to reveal $\mathbf{k}_{w,0}^i$ to P_i .
 - (c) P_i sets the one key as $\mathbf{k}_{w,1}^i \leftarrow \mathbf{k}_{w,0}^i \oplus R^i$ and all parties set $\llbracket \mathbf{k}_{w,1}^i \rrbracket_{2^k} \leftarrow \llbracket \mathbf{k}_{w,0}^i \rrbracket_{2^k} \oplus \llbracket R^i \rrbracket_{2^k}$.
 - (d) For all four distinct values of $(\alpha, \beta) \in \{0, 1\}^2$, and for every $j \in [n]$, P_i calls $\mathcal{F}_{\text{Prep}}$ with input $(\text{Input}, i, \llbracket F_{\alpha,\beta}^{g,i,j} \rrbracket_{2^k}, F_{\mathbf{k}_{u,\alpha}^i, \mathbf{k}_{v,\beta}^i}(g\|j), 1)$ and the other parties with input $(\text{Input}, i, \llbracket F_{\alpha,\beta}^{g,i,j} \rrbracket_{2^k}, \perp, 1)$.
 3. For all $j \in [n]$ and all $(\alpha, \beta) \in \{0, 1\}^2$, the parties compute

$$\begin{aligned} \llbracket \tilde{g}_{\alpha,\beta}^j \rrbracket_{2^k} \leftarrow & \left(\bigoplus_{i=1}^n \llbracket F_{\alpha,\beta}^{g,i,j} \rrbracket_{2^k} \right) \oplus \llbracket \mathbf{k}_{w,0}^j \rrbracket_{2^k} \\ & \oplus \llbracket R^j \rrbracket_{2^k} \cdot ((\llbracket \lambda_u \rrbracket_{2^k} \oplus \alpha) \cdot (\llbracket \lambda_v \rrbracket_{2^k} \oplus \beta) \oplus \llbracket \lambda_w \rrbracket_{2^k}) \end{aligned}$$

Fig. 18. Subprotocol $\Pi_{\text{ABB+BMR}}^{\text{Garble}}$

Subprotocol $\Pi_{\text{ABB+BMR}}^{\text{Eval}}$

Evaluate The parties, holding the output $((\tilde{g}_{0,0}^i, \tilde{g}_{0,1}^i, \tilde{g}_{1,0}^i, \tilde{g}_{1,1}^i)_{i=1}^n)_{g \in G}$ of $\mathcal{F}_{\text{BMR}}^{\text{Garble}+}$, evaluate in the following way, traversing the circuit in topological order:

1. For each input $\{\llbracket a_i \rrbracket_p\}_{i \in [t]}$, the parties do the following:
 - (a) Retrieve from memory the secret mask $\llbracket r_i \rrbracket_p$ produced in **Input layer**.
 - (b) Compute the secret $\llbracket x_i \rrbracket_p \leftarrow \llbracket a_i \rrbracket_p - \llbracket r_i \rrbracket_p$.
 - (c) Call $\mathcal{F}_{\text{Prep}}$ with input $(\text{Open}, \llbracket x_i \rrbracket_p, 0)$.
 - (d) Denote the corresponding input wires by $\{w_{i,j}\}_{j=0}^{\lceil \log p \rceil - 1}$. Bit-decompose the public value x_i and let the bits be $\{x_{i,j}\}_{j=0}^{\lceil \log p \rceil - 1}$.
 - (e) For each $j = 0, \dots, \lceil \log p \rceil - 1$, retrieve from memory the wire mask $\lambda_{w_{i,j}}$ from **Input layer** for $x_{i,j}$ and set $\Lambda_{w_{i,j}} \leftarrow x_{i,j} \oplus \lambda_{w_{i,j}}$.
 - (f) For each $l \in [n]$, P_l sends $\{k_{w_{i,j}, \Lambda_{w_{i,j}}}^l\}_{j=0}^{\lceil \log p \rceil - 1}$ to all other parties.
2. For every $g \in G$,
 - (a) If g is an XOR gate,
 - i. Party P_i computes $\Lambda_w \leftarrow \Lambda_u \oplus \Lambda_v$.
 - ii. Party P_i computes all n output keys indexed by $j \in [n]$, as $k_{w, \Lambda_w}^j \leftarrow k_{u, \Lambda_u}^j \oplus k_{v, \Lambda_v}^j$.
 - (b) If g is an AND gate,
 - i. Each party computes the n keys indexed by $j \in [n]$ as

$$k_{w, \Lambda_w}^j \leftarrow \tilde{g}_{\Lambda_u, \Lambda_v}^j \oplus \left(\bigoplus_{i=1}^n F_{k_{u, \Lambda_u}^i, k_{v, \Lambda_v}^i}(g \parallel j) \right)$$

and compares its keys $k_{w,0}^i$ and $k_{w,1}^i$ to the i^{th} key obtained to determine the global signal bit Λ_w .

3. For every external output wire w ,
 - (a) Retrieve from memory the corresponding public signal bit Λ_{w_0} produced in **Output layer**.
 - (b) Locally compute $\Lambda_{w'} \leftarrow \Lambda_{w_0} \oplus \Lambda_w$.
 - (c) Locally compute the secret output as

$$\llbracket b_w \rrbracket_p \leftarrow \Lambda_{w'} + \llbracket \lambda_{w'} \rrbracket_p - 2 \cdot \Lambda_{w'} \cdot \llbracket \lambda_{w'} \rrbracket_p.$$

4. Send the message $(\text{Check}, 1)$ to $\mathcal{F}_{\text{Prep}}$.

Fig. 19. Subprotocol $\Pi_{\text{ABB+BMR}}^{\text{Eval}}$

D SPDZ-BMR PRF Assumption

The non-existence of *circular 2-correlation robust* PRFs required for using the multiparty FreeXOR technique would force garbling protocols to garble XOR gates in the same way

as AND gates, providing PRFs under the (supposed) weaker assumption of *pseudorandom function under multiple keys* exist. These are defined as follows:

Definition D.1. Let $F : \{0, 1\}^\kappa \times \{0, 1\}^\kappa \rightarrow \{0, 1\}^\kappa$ be an efficient, length-preserving, keyed function. We say that F is a pseudorandom function under multiple keys if for all polynomial time distinguishers \mathcal{D} there exists a negligible function ν such that:

$$\left| \Pr[\mathcal{D}^{F_{\mathbf{k}}}(\mathbf{1}^{\text{sec}}) = 1] - \Pr[\mathcal{D}^{\bar{f}}(\mathbf{1}^{\text{sec}}) = 1] \right| \leq \nu(\kappa).$$

where $F_{\mathbf{k}}$ denotes the tuple $(F_{\mathbf{k}_1}, \dots, F_{\mathbf{k}_n})$ of the pseudorandom function F keyed using $\mathbf{k}_1, \dots, \mathbf{k}_n$ and \bar{f} denotes the tuple (f_1, \dots, f_n) of random functions $\{f_i : \{0, 1\}^\kappa \rightarrow \{0, 1\}^\kappa\}_{i=1}^n$.

Lindell et al. proved that the SPDZ-BMR technique is secure under this assumption on the PRF.

E MASCOT and LowGear

Here we provide more data on the communication complexity of these two protocols benchmarked for fields \mathbb{F}_p and \mathbb{F}_{2^k} where $k = 128$, $\log p > 128$, and statistical security $\text{sec} = 64$.

# Parties	MASCOT \mathbb{F}_{2^k}		LowGear \mathbb{F}_p	
	Input (bit)	Triple	Input	Triple
2	0.384	360.44	2.048	30.146
3	1.024	1081.32	5.888	89.67
4	1.92	2162.64	11.520	178.572
5	3.072	3604.4	18.94	296.85

Table 6. Communication costs (kbits) for fields with different characteristic.